

**2DV605**

# **Parallel Computing**

## **Architectures, Models, Performance**

**Sabri Pllana, PhD**

**Department of Computer Science**

**Building D, Room D2236C, LNU, Växjö**

**sabri.pllana@lnu.se, <http://homepage.lnu.se/staff/saplaa/>**



# Outline

- ☐ **Parallel Computer Architectures**
- ☐ **Parallel Programming Models**
- ☐ **Performance**
- ☐ **Presentation Techniques**



# Selected Literature

## ❑ Textbooks

- J. Hennessy and D. Patterson. “[Computer Architecture: A Quantitative Approach](#)”. Morgan Kaufmann, 5th Edition, September 2011.
- R. Jain. “[The Art of Computer Systems Performance Analysis: Techniques for Experimental Design, Measurement, Simulation, and Modeling](#)”. Wiley- Interscience, New York, NY, April 1991.
- S. Pillana and F. Xhafa. “[Programming multi-core and many-core computing systems](#)”, John Wiley & Sons, Inc., Hoboken, NJ, USA, 2017.

## ❑ Performance Evaluation: Projects and Tools

- Paradyn Tools Project ([www.cs.wisc.edu/paradyn/](http://www.cs.wisc.edu/paradyn/))
- TAU ([www.cs.uoregon.edu/research/tau/](http://www.cs.uoregon.edu/research/tau/))
- VI-HPS (<http://www.vi-hps.org/tools/>)



# Parallel Computing Architectures



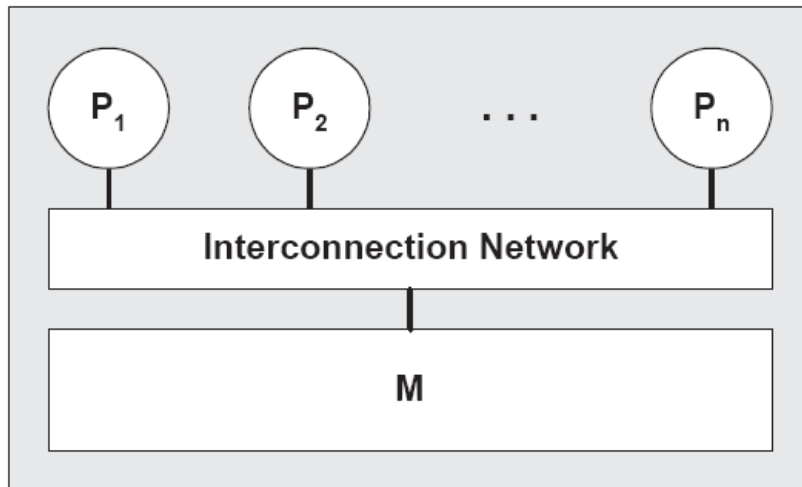
# Shared and Distributed Memory Systems

## ❑ Shared memory (SM)

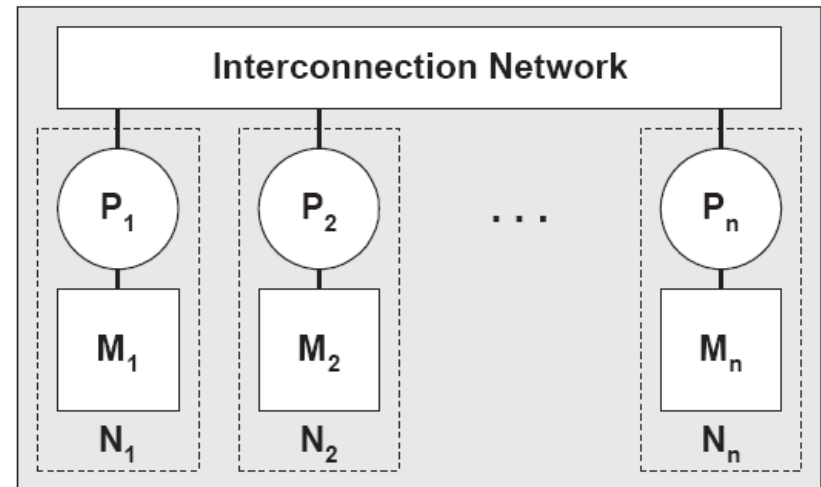
- **easy to program**, but do not scale to a large number of processors

## ❑ Distributed memory (DM)

- **scalable**, but harder to program than SM systems



*Shared memory*

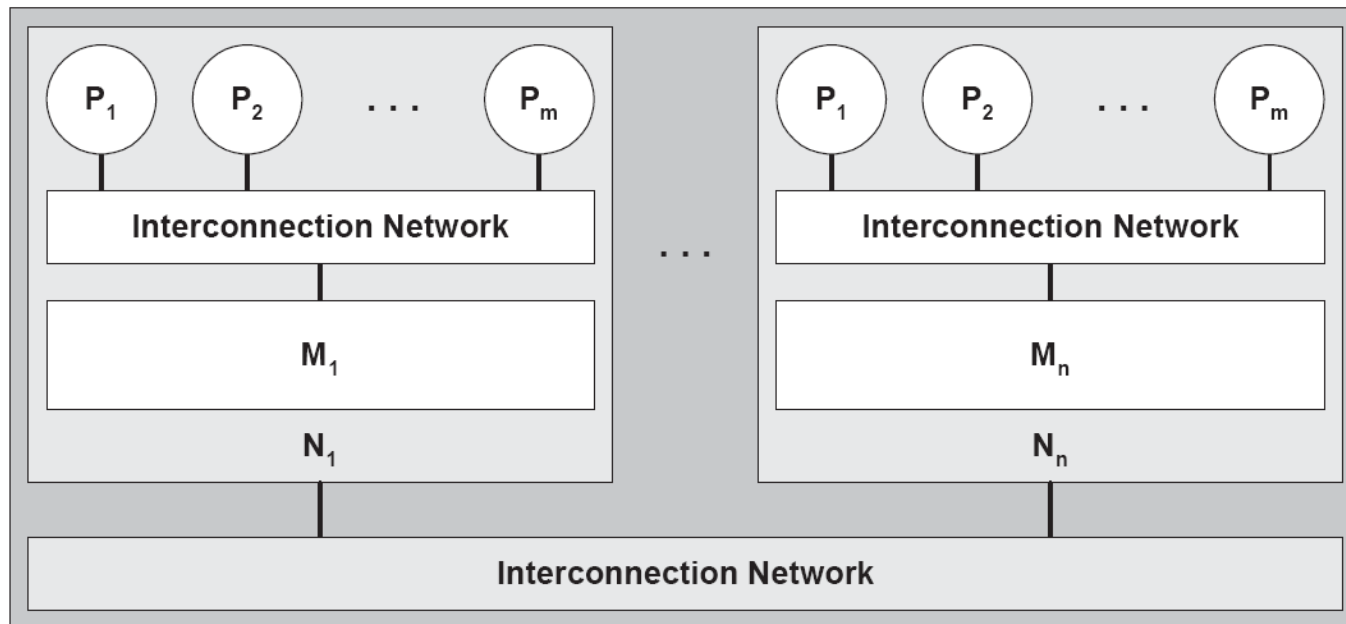


*Distributed memory*



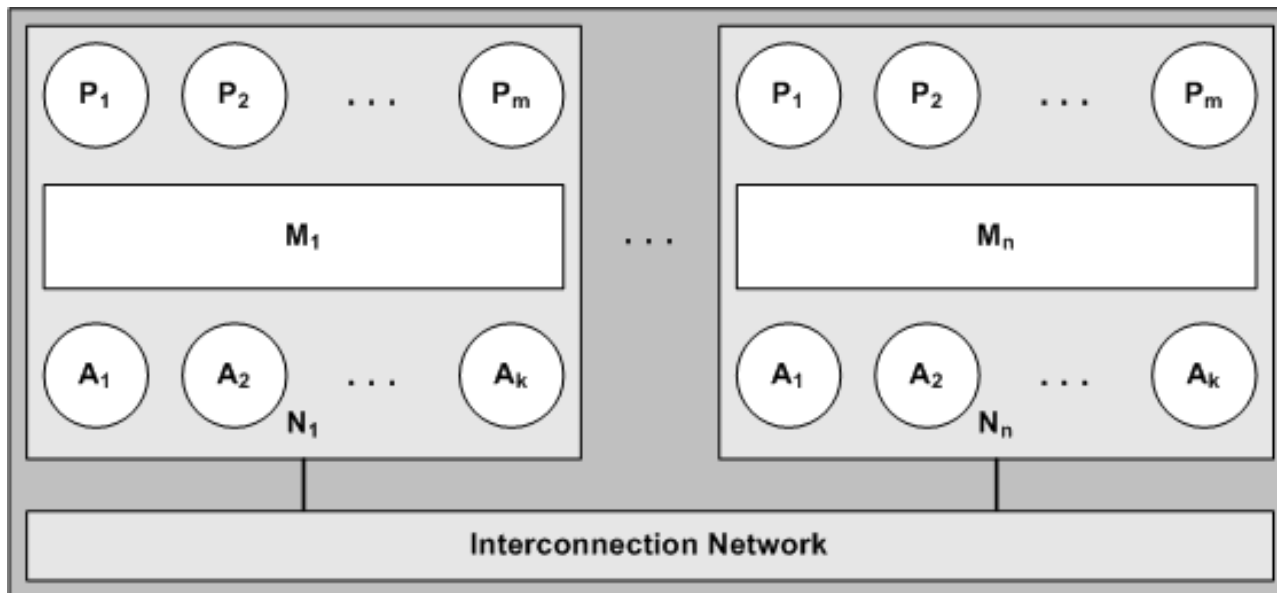
# Clusters of SMPs

- ❑ **Clusters of symmetric multiprocessing (SMP) systems**
  - combine features of SM and DM systems



# Accelerated Systems

- ❑ Also known as *heterogeneous* or *asymmetric systems*
- ❑ Common: GPU-accelerated systems



# GPU Example: NVIDIA Volta Architecture





- V100: 7.5 TF of double-precision and 15 TF of single-precision performance
- 80 Streaming Multiprocessors (SMs)
- an SM has 64 FP32 cores, 64 INT32 cores, 32 FP64 cores, 8 tensor cores, 4 texture units
- NVLink enables communication among CPU and GPUs; 5 to 12x faster than the third generation of PCIe





# Cerebras Wafer-Scale Engine (WSE)

- *optimized for deep learning*
- *400 000 cores*
- *18 GB on-chip SRAM*
- *WSE = 56x GPU size*
- *100 – 1000 faster than NVIDIA V100 GPU*



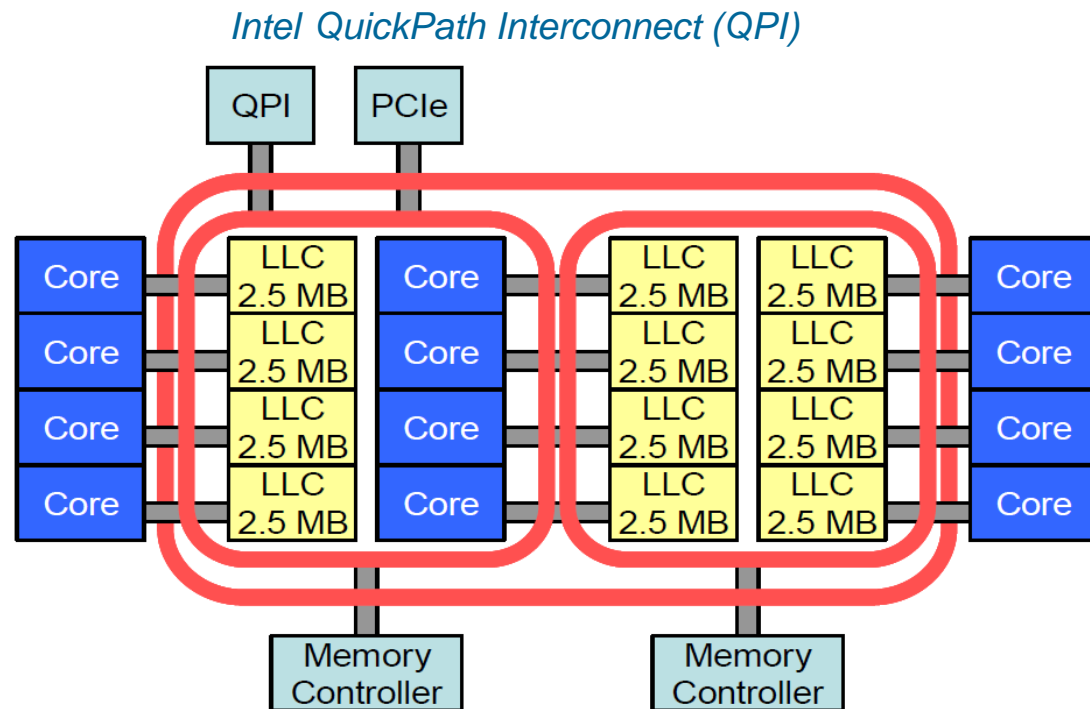
WSE

GPU



# Intel Xeon Processor E5-2600 v2

- *Intel Xeon E5-2695v2 comprises 12 cores*
- *up to 30 MB of last level cache*



*Credit: IBM Systems and Technology Group*



# Examples: Old Clusters at the University of Vienna



*Cluster of SMPs*




*GPU-accelerated System*





# LNU Parallel Computing Lab: Ida

Ida		
Specs	Intel Xeon E5	GeForce GPU
Type	E5-2650 v4	GTX Titan X
Core Frequency	2.2 - 2.9 GHz	1 - 1.1 GHz
# of Cores	12	3072
# of Threads	24	/
Cache	30 MB	/
Mem. Bandwidth	76.8 GB/s	336.5 GB/s
Memory	384 GB	12 GB
TDP	105 W	250 W



# LNU Parallel Computing Lab: DISA

- *Installed at LNU in 2018*
- *19 DELL EMC nodes:*
  - 14 PowerEdge R740
  - 5 PowerEdge R740XD
- *Various roles of nodes:*
  - homogeneous computing,
  - heterogeneous computing,
  - data storage,
  - system administration,
- *Heterogeneous nodes contain two CPUs and four GPUs*

Specs	Intel Xeon	NVIDIA GPU
Type	Gold 6148	Quadro P4000
Frequency (GHz)	2.4 - 3.7	1.48
# of Cores	20	1792
# of Threads	40	/
Cache (MB)	27.5	/
Memory (GB)	768	8
TDP (W)	150	105



Front



Inside a Node



Back

# Examples: Summit at ORNL

- *Rank 1 in TOP500 list (June 2018)*
- *200 petaflops (max performance)*
- *250 petabytes of data storage*
  - *74 years of HD video*
- *4608 computing nodes*
- *each node comprises*
  - *two IBM Power9 22-core CPUs*
  - *six NVIDIA Volta GPUs*
- *340 tons*
- *space of two tennis courts*



# Parallel Programming Models






# Software for Parallel Computing Systems

To exploit the performance of parallel computing systems, application programs have to be *parallelized*


## □ Parallelization

- Data distribution
- Work distribution
- Communication/Synchronization




## □ Parallel APIs and Languages

- MPI, OpenMP, OpenCL, ...



## □ Libraries

- LAPACK, ScaLAPACK, ...



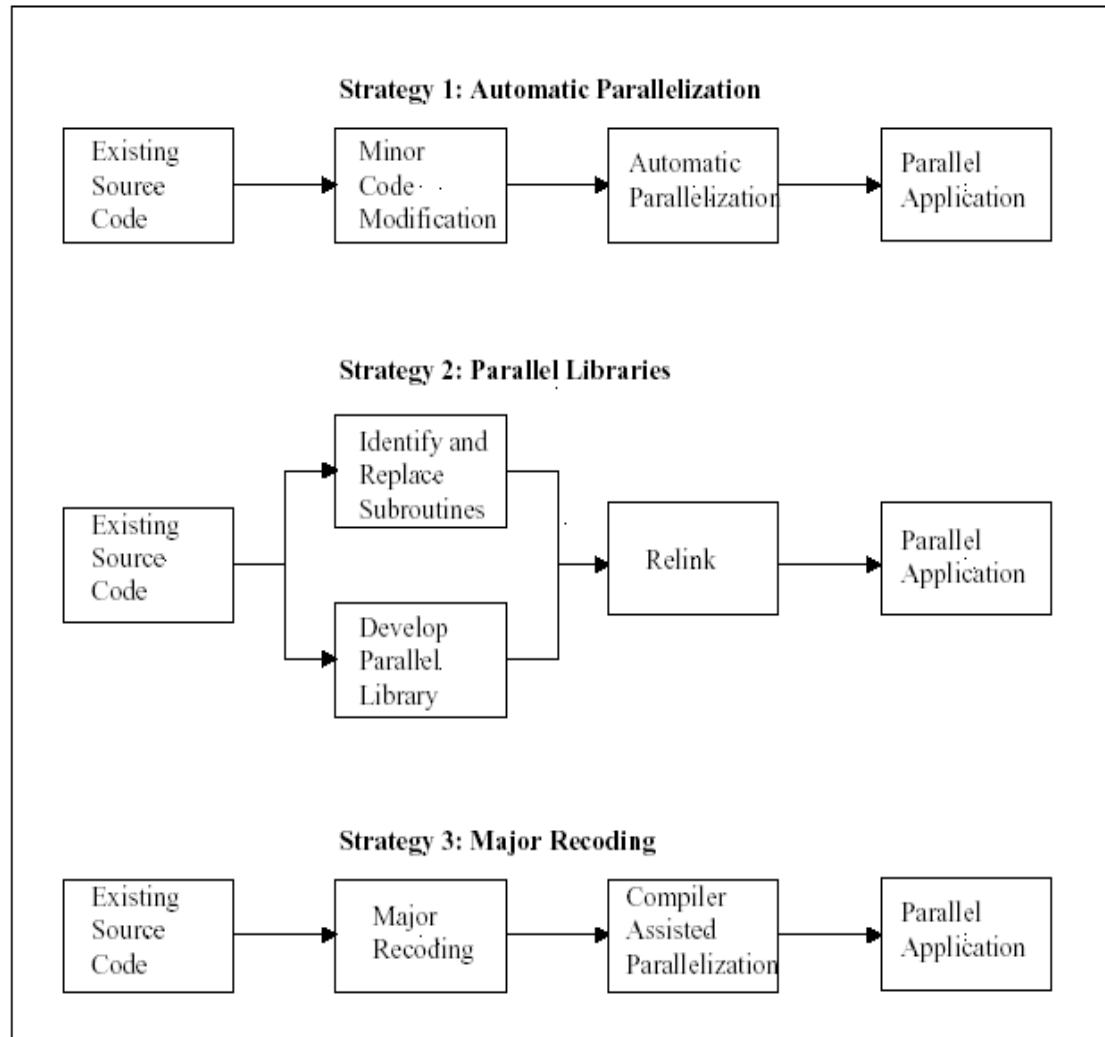
## □ Programming parallel computing systems is complex

- *“no free lunch”*





# Parallelization (1)



# Parallelization (2)

## ❑ Focus on two major resources

- processors
- memory

## ❑ Parallelization involves

- distributing work to processors
- distributing data (if memory is distributed)

and

- synchronization of the distributed work
- communication of remote data to local processor (if memory is distr.)

❑ **Programming models** offer a combined method for distribution of work & data, synchronization and communication

- Distributed-Memory Programming Models
- Shared-Memory Programming Models



# Programming Models

## ☐ Distributed Memory Model (Message Passing Model)

- MPI
- Applicable also to shared-memory machines

## ☐ Shared Memory Model

- OpenMP, pThreads, Java threads,
- Restricted to shared-memory machines

## ☐ Hybrid Model

- Mixing shared and distributed memory model
- Using OpenMP and MPI together

## ☐ Other Programming Models

- Accelerated or heterogeneous systems: CUDA, OpenCL, OpenACC
- Partitioned Global Address Space (PGAS): Chapel, X10

## ☐ Object and Service Oriented Models

- OO: RPC/RMI
- Web Services (Grid or Cloud computing)



# Work & Data Distribution

## □ Work distribution

- usually based on loop decomposition

do i=1,100

→ i=1,25

i=26,50

i=51,75

i=76,100

## □ Data distribution

- all work for a local portion of the data is done by the local processor

A( 1:20, 1: 50)



A( 1:20, 51:100)

A(21:40, 1: 50)

A(21:40, 51:100)

## □ Domain decomposition

- decomposition of work and data is done in a higher model, reflecting reality



# Single-Program Multiple-Data (SPMD)

## Domain Decomposition

1. Divide domain (=data) into different parts.
2. Assign parts to different processors.


## Owner-Computes Rule

- Each data point is owned by some processor
- The owner of a data point is responsible for computing its value.
- Communication (explicitly or implicitly) is performed, if computation of a data point requires data from other processors.

Domain decomposition typically leads to SPMD structure.

## Single-Program Multiple-Data (SPMD)

- Each processor executes the same code, but on different data
- Scalar variables are typically replicated on all processors and redundantly computed on each of them.



# Communication

```
Do i=2,99  
  b(i) = a(i) + f*(a(i-1)+a(i+1)-2*a(i))  
Enddo
```

Communication is necessary on the boundaries

– e.g.  $b(26) = a(26) + f*(a(25) + a(27) - 2*a(26))$


**A** (1 : 25)

**A** (26 : 50)

**A** (51 : 75)

**A** (76 : 100)


– e.g. at domain boundaries



# Synchronization

## Synchronization

- is needed in some contexts
- may cause
  - idle time on some processors
  - overhead to execute the synchronization primitive



# Dependence

- ❑ Dependence is a relationship that places constraints on the *execution order of computations*

- determines whether or not computations can be executed in parallel.

- ❑ There are two types of dependences

- Control dependence

```
S1: if (b ≠ 0) then  
S2:   c = a / b;
```

- Data dependence, relation on the set of statements

```
S1: a = ...  
S2: c = a / b;
```





# MPI and OpenMP

## ❑ MPI (Message Passing Interface)

- user specifies how work & data is distributed
- user specifies how and when communication/synchronization is done
- by calling MPI communication library-routines
- applicable to distributed- and shared-memory architectures

## ❑ OpenMP

- shared memory directives to define the work decomposition
- high-level multi-threading
- restricted to shared-memory systems
- synchronization is implicit (can be also user-defined)




# Performance



# System Performance

- ❑ Performance =  $f(\text{system}, \text{workload})$
- ❑ In our case:
  - workload  $\sim$  computer program
  - system  $\sim$  computer system
  - performance  $\sim$  [program execution time](#), power consumption



# Performance of Computing Systems

- ❑ **Expressed in floating-point operations per second (FLOPS)**
- ❑ **Peak performance**
  - theoretical maximal number of FLOPS
  - all components of the system are completely utilized
- ❑ **Sustained Performance**
  - obtained FLOPS when a certain program is executed



# Performance of Processor

## ❑ Clock frequency of processor

- speed at which are performed the fundamental operations
- for instance: add operation
- clock frequency is expressed in **cycles per second**
- one cycle per second corresponds to one **hertz (Hz)**
- 1GHz: 1 000 000 000 (or  $10^9$ ) cycles per second
- clock frequency is also known as the **speed of processor**

## ❑ Performance depends not only on clock frequency

- architecture of processor is important



# Program Execution Time

- ❑ Clock frequency:  $F$
- ❑ Number of clock cycles:  $NC$
- ❑ Number of instructions:  $NI$
- ❑ Clocks per instruction:  $CPI = NC / NI$
  
- ❑ Program execution time =  $NI * CPI * 1/F$ 
  - also known as Iron law
  
- ❑ Average time per instruction:  $TI_{av} = CPI * 1/F$ 
  - program execution time =  $NI * TI_{av}$



# Memory Performance

- ❑ **Data transfer between the memory and processor**
- ❑ **Latency (also known as delay)**
  - time between the request and the retrieve of a data element from the memory
- ❑ **Bandwidth**
  - quantity of data that can be transferred between the processor and the memory within one second



# Network Performance

- ❑ **Data transfer among the nodes via the interconnection network**
- ❑ **Latency (also known as delay)**
  - time that is needed to transfer a data element between two nodes of network
- ❑ **Bandwidth**
  - number of bits of data that can be transferred between two nodes within one second





# Speedup

- ❑ Number of processors:  $NP$
- ❑ Program execution time on one processor:  $T_1$
- ❑ Program execution time on  $NP$  processors:  $T_{NP}$
- ❑ Speedup:  $Speedup(NP) = \frac{T_1}{T_{NP}}$
- ❑ Linear speedup:  $Speedup(NP) = NP$
- ❑ Super linear:  $Speedup(NP) > NP$



# Efficiency

- ❑ Number of processors: **NP**

$$Efficiency(NP) = \frac{Speedup(NP)}{NP}$$



# Scalability

- ❑ Performance improvement is proportional to the increase of the number of processors
- ❑ Theoretical maximal speedup can be estimated using **Amdahl's law**
  - fraction of computation that is performed sequentially:  $\sigma$
  - fraction of computation that is performed in parallel:  $1 - \sigma$

$$Speedup(\sigma, NP) = \frac{1}{\sigma + \frac{1-\sigma}{NP}}$$


- maximal speedup depends only on the fraction of computation that is performed sequentially

$$\lim_{NP \rightarrow \infty} \frac{1}{\sigma + \frac{1-\sigma}{NP}} = \frac{1}{\sigma}$$



# Speedup and Amdahl's law

*“the effort expended on achieving high parallel processing rates is wasted unless it is accompanied by achievements in sequential processing rates of very nearly the same magnitude”* - Gene M. Amdahl, 1967



$$\text{Speedup} = T_{\text{original program}} / T_{\text{improved program}}$$



# Performance Evaluation Methods

## ❑ Measurement

- **trustfulness**: results obtained on a **real** system
- requires that the system under study is available

## ❑ Mathematical modeling

- represents the system as a set of symbolic expressions
- lacks the structural information (type/relationship of SW/HW )

## ❑ Simulation

- imitates the behavior of system under study
- commonly implemented as a computer program
- may be time-consuming



# Benchmarks

- ❑ **Benchmarks are computer programs**
  - artificial (instruction mixes) or real programs
  - measure various aspects (computation, communication, power)
- ❑ **LINPACK benchmark**
  - solves dense systems of linear equations
  - is used to rank 500 most powerful computing systems
  - TOP500 is updated two times per year ([www.top500.org](http://www.top500.org))
- ❑ **Standard Performance Evaluation Corporation (SPEC)**
  - SPEC benchmarks: CPU, MPI/OMP, Power, ...
  - [www.spec.org](http://www.spec.org)



# June 2019

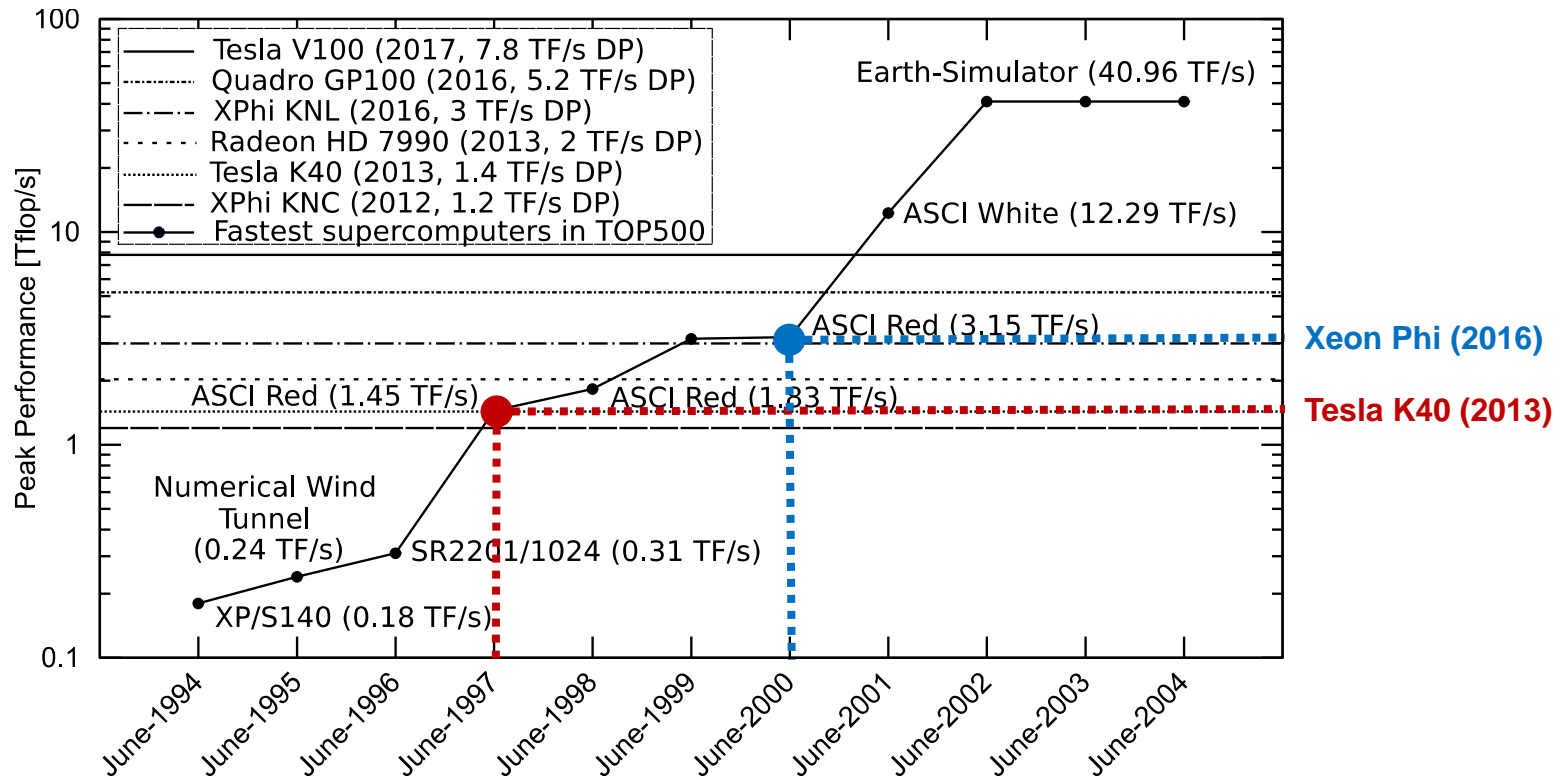
## www.top500.org

- **Rmax** is maximal performance achieved in the High Performance Computing LINPACK benchmark
- **Pflop/s** or petaflop:  $10^{15}$  FLOPS

		SPECS	SITE	COUNTRY	CORES	R <sub>MAX</sub> PFLOP/S	POWER MW
1	Summit	IBM POWER9 (22C, 3.07GHz), NVIDIA Volta GV100 (80C), Dual-Rail Mellanox EDR Infiniband	DOE/SC/ORNL	USA	2,414,592	148.6	11.4
2	Sierra	IBM POWER9 (22C, 3.16GHz), NVIDIA Tesla V100 (80C), Dual-Rail Mellanox EDR Infiniband	DOE/NNSA/LLNL	USA	1,572,480	94.6	7.44
3	Sunway TaihuLight	Shenwei SW26010 (260C, 1.45 GHz) Custom Interconnect	NSCC in Wuxi	China	10,649,600	93.0	15.4
4	Tianhe-2A (Milkyway-2A)	Intel Ivy Bridge (12C, 2.2 GHz) & TH Express-2, Matrix-2000	NSCC Guangzhou	China	4,981,760	61.4	18.5
5	Frontera	Dell C6420, Xeon Platinum 8280 28C 2.7GHz, Mellanox InfiniBand HDR	TACC/U of Texas	USA	448,448	23.5	-



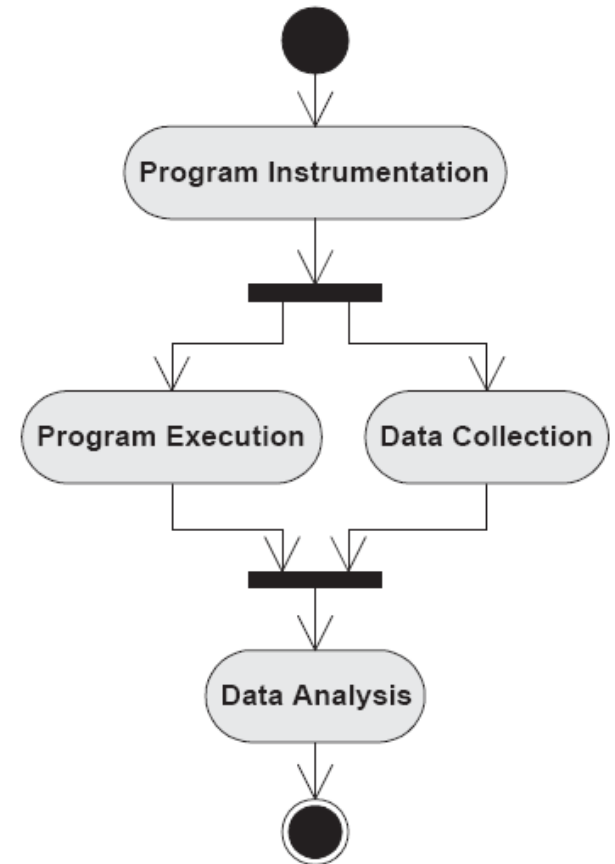
# Performance of Accelerators and #1 in TOP500





# Program Performance Measurement

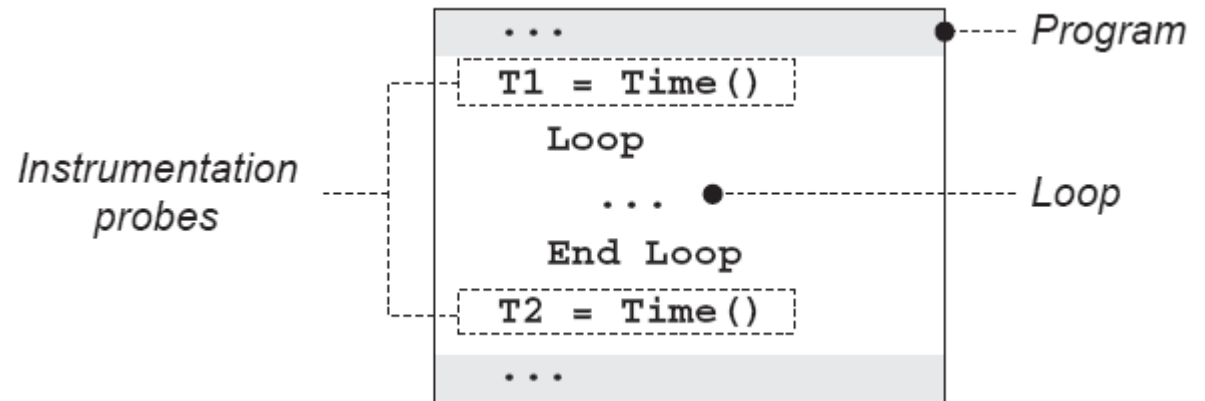
- ❑ **Common activities**
  - instrumentation
  - execution
  - data collection
  - analysis



# Program Instrumentation

## ❑ Instrumentation probes

- program statements
- read and store the current time
- may be done manually or automatically
- loop execution time:  $T_{Loop} = T_2 - T_1$



# Data Collection

## ❑ Examples of the collected data

- CPU time
- number of floating point instructions
- number of times a function is executed
- number of load/store instructions

## ❑ Data sources

- instrumented program
- operating system (CPU utilization)
- hardware performance counters (L2 requests)



# Data Analysis: Average Value

## ❑ Multiple program executions

- same program on the same machine
- may produce various performance results
- **outliers**: data values that differ significantly from other values
- outliers are ignored if there is no reasonable explanation

## ❑ Mean **M** of a set of measurements $\{x_1, .., x_N\}$

$$M = 1/N \sum_{i=1}^N x_i$$



# Data Analysis: Variance and Standard Deviation

## □ Variance

$$S^2 = 1/(N - 1) \sum_{i=1}^N (x_i - M)^2$$

## □ Standard deviation

$$S = \sqrt{1/(N - 1) \sum_{i=1}^N (x_i - M)^2}$$



# Data Analysis: Standard Error

- ❑ **Various sets of measurements**
  - may result with various means  $M$
- ❑ **Standard error**
  - is the standard deviation of the mean  $M$
- ❑ **Standard error of a set of measurements  $\{x_1, \dots, x_N\}$**

$$S_M = S/\sqrt{N}$$



# Data Analysis: Systematic Error

## ☐ Biased measurements

- are reproducible
- are always biased in the same direction (too large or too small)

## ☐ Possible sources of systematic errors

- instrumentation issues: overhead, imperfection,..
- methodological mistakes

## ☐ Systematic errors are hard to detect and correct

- can not be corrected by simply repeating the experiment
- may be helpful [to try various measurement tools and methods](#) or use [calibration](#)



# Data Analysis: Confidence Interval (1)

## ❑ Assume normal distribution of measurement results

- mean:  $\mu$
- standard deviation:  $\sigma$

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

## ❑ Confidence interval of $k\sigma$ is calculated as follows

$$P(\mu - k\sigma < x < \mu + k\sigma) = \frac{1}{\sigma\sqrt{2\pi}} \int_{\mu-k\sigma}^{\mu+k\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

## ❑ Commonly used intervals

- 95% corresponds to  $1.96\sigma$
- 99% corresponds to  $2.58\sigma$
- higher levels of confidence imply larger confidence intervals






# Data Analysis: Confidence Interval (2)

## □ Data analysis

- repeat experiments to deal with **performance fluctuation**
- average value, outliers, variance, standard deviation and error


## □ Confidence intervals

- $\mu \pm \sigma$ , 68.3% of observations
- $\mu \pm 2\sigma$ , 95.4%
- $\mu \pm 3\sigma$ , 99.7%



# Normal Distribution (or Gaussian Distribution)

- Carl Friedrich Gauß (1777 – 1855)
- central limit theorem: *the **mean** of a large set of independent random variates is distributed approximately normally*



*Math is everywhere!*



# EPCC OpenMP Benchmark

- ❑ Comparison of OpenMP implementations
- ❑ EPCC: Edinburgh Parallel Computing Center
- ❑ Estimation of OMP overheads:  $T_{ovr}$ 
  - parallel execution:  $T_{par}$
  - sequential execution:  $T_{seq}$
  - number of processors:  $NP$

$$T_{ovr} = T_{par} - \frac{T_{seq}}{NP}$$



# EPCC OpenMP Benchmark: Example

## ❑ PARALLEL DO overhead

- specifies that iterations of the loop should be executed in parallel

nthreads = 4

```
1:      do k = 0, outerreps
2:          start = getclock()
3:          do j = 1, innerreps
4:      !$OMP PARALLEL DO
5:          do i = 1, nthreads
6:              call delay(dl)
7:          enddo
8:      !$OMP END PARALLEL DO
9:          end do
10:         time(k) = (getclock() - start) * 1.0e6 / dble (innerreps)
11:     end do
```

overhead of  
PARALLEL DO is  
 $5.16 \pm 0.21$   
microseconds

### Code example

Sample_size	Average	Min	Max	S.D.	Outliers
50	11.10000	11.00000	11.20000	0.05714	0
PARALLEL_DO time =			11.10 microseconds +/- 0.112		
PARALLEL_DO overhead =			5.16 microseconds +/- 0.210		

### Result




# SKaMPI Benchmark


- ❑ **Assessment of various MPI implementations**
- ❑ **SKaMPI: Special Karlsruher MPI Benchmark**
- ❑ **Support**
  - point-to-point communication
  - collective communications



# SKaMPI Benchmark: MPI Send, MPI Receive



*execution time of MPI SEND operation for various message lengths*




*time needed to SEND and RECEIVE messages of various lengths*

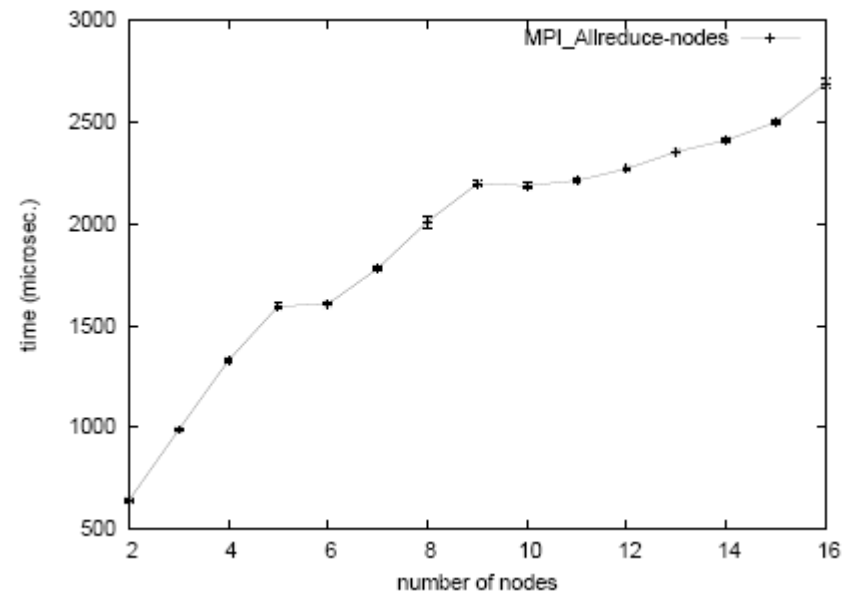


# SKaMPI Benchmark: MPI Allreduce

performs a specific *operation* (such as *sum*, *min*, *max*) on supplied *data*, and broadcasts the result back to all involved processes



*MPI Allreduce for various message lengths*



*MPI Allreduce for various number of nodes*



# Presentation Techniques





# Presentation of Results

- ❑ **We provide several recommendations**
  - based on the Chapter 10 of the book "*The Art of Computer Systems Performance Analysis*" of Raj Jain
- ❑ **We expect that you will use these techniques for presenting your results**
  - quality and completeness of your presentation will be assessed
- ❑ **Worth to read**
  - "*Twelve Ways to Fool the Masses When Giving Performance Results on Parallel Computers*" by David H. Bailey
  - Slides: <http://www.davidhbailey.com/dhbtalks/dhb-12ways.pdf>
  - Paper: <http://crd-legacy.lbl.gov/~dhbailey/dhbpapers/twelve-ways.pdf>



# Graphic Charts

- ❑ **Line chart**
  - x-axis: **cause** or independent variable
  - y-axis: **effect** or dependent variable
- ❑ **Bar chart**
- ❑ **Pie chart**
- ❑ **Histograms**
  - visualizes distribution of data
- ❑ **Gantt chart**
  - may be used to visualize activities of parallel processes
- ❑ **Kiviat chart**
  - two-dimensional representation of more than two variables



# Variables

## ☐ Qualitative

- **ordered**: type of computers (notebook, supercomputer)
- **unordered**: types of programs (entertainment, education)

## ☐ Quantitative

- **discrete**: integer number (number of processes or threads)
- **continuous**: real number (execution time)

## ☐ Type of variable determines the type of graphic chart

- line chart is used when **x** and **y** are continuous variables
- bar chart is suitable for discrete **x** variables



# Guidelines for Graphic Charts

- ❑ **Minimize the reading effort**
  - should be easy to understand the chart
- ❑ **Maximize the conveyed information**
  - use understandable key words; [label properly axes](#)
- ❑ **Minimize ink**
  - consider removing grid lines from the graph
- ❑ **Use common practices**
  - [show dependent variables using y-axis](#)
- ❑ **Avoid ambiguity**
  - identify/label individual elements of the chart



# Checklist for Graphical Presentation of Results

- ❑ **A selection from Page 143, Chapter 10**
  - Are both axes labeled
  - Are the labels self-explanatory
  - Are the scales and divisions shown on both axes
  - Are the curves on a line chart individually labeled
  - Are the units of measurement indicated
  - Is the horizontal scale increasing from left to right
  - Is the vertical scale increasing from bottom to top
  - Are the grid lines aiding in reading the curves
  - Does this whole chart add to information available to the reader




# Common Mistakes in Preparing Graphic Charts


- ❑ **Presenting too many alternatives on a single chart**
  - avoid too many curves or bars ("*small is beautiful*")
- ❑ **Using symbols in place of text**
  - the reader should not be forced to decipher the meaning of symbols to understand the chart
- ❑ **Placing extraneous information on the chart**
  - avoid any detracting elements (grid lines)
- ❑ **Selecting scale ranges improperly**
  - select appropriate minimum and maximum values
- ❑ **Using not the appropriate type of chart**
  - avoid using a line chart instead of bar chart



# Common Mistakes: Example



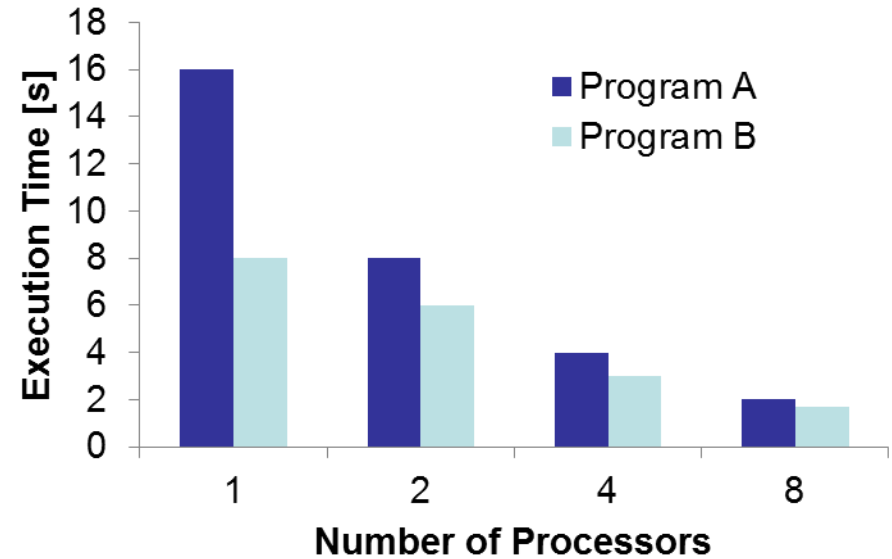

# Showing Only Speedup Results is not Enough






# Speedup and Execution Time Results

Which program is better?



# Describe the Experimentation Platform



When we report performance results, we usually specify the experimentation platform: program, OS, HW, RT software, compiler type and parameters



# Selected Linux Commands

- ❑ Environment information
  - `env`
- ❑ Operating system information
  - `uname -a`
- ❑ CPU information
  - `cat /proc/cpuinfo`
- ❑ Memory information
  - `cat /proc/meminfo` or
  - `free -m`



# Command “time”

## ❑ A UNIX command

- measures program execution time

## ❑ Displays the following information

- **user** time (time CPUs spent for executing the program in user space)
- **system** time (time CPUs spent for system calls in kernel mode)
- **real** time (or wall clock time) between program invocation and termination

## ❑ Example

- `time mpirun -np 4 matvec_1000`
- `0.32u 0.10s 0:00.19`

Why the user time plus system time is larger than the real time?



# Command “perf”

## ❑ A Linux performance measurement program

- example: `perf stat -e cache-misses ./matmul`
- example: `perf stat ./matmul`

```
-bash-4.1$ perf stat ./matmul
Execution time: 0.164323 [s] // output of matmul program
```

Performance counter stats for './matmul':

6222.036262	task-clock	#	36.071	CPUs utilized
188	context-switches	#	0.030	K/sec
60	cpu-migrations	#	0.010	K/sec
1,074	page-faults	#	0.173	K/sec
7,459,088,185	cycles	#	1.199	GHz
4,647,863,937	stalled-cycles-frontend	#	62.31%	frontend cycles idle
<not supported>	stalled-cycles-backend			
6,683,693,162	instructions	#	0.90	insns per cycle
		#	0.70	stalled cycles per insn
672,748,725	branches	#	108.124	M/sec
292,465	branch-misses	#	0.04%	of all branches

0.172492173 seconds time elapsed

```
-bash-4.1$ perf stat -e cache-misses ./matmul
Execution time: 0.146337 [s] // matmul output
```

Performance counter stats for './matmul':

90,633 cache-misses

0.155192113 seconds time elapsed



# Tuning and Analysis Utilities (TAU)

## ❑ Support


- code instrumentation
- [profiling](#) (distribution of execution time across routines)
- tracing (shows when and where an event occurred)
- performance analyzing (textual and graphical display of results)

## ❑ Documentation is available online

- <http://www.cs.uoregon.edu/research/tau/docs.php>



# TAU: ParaProf

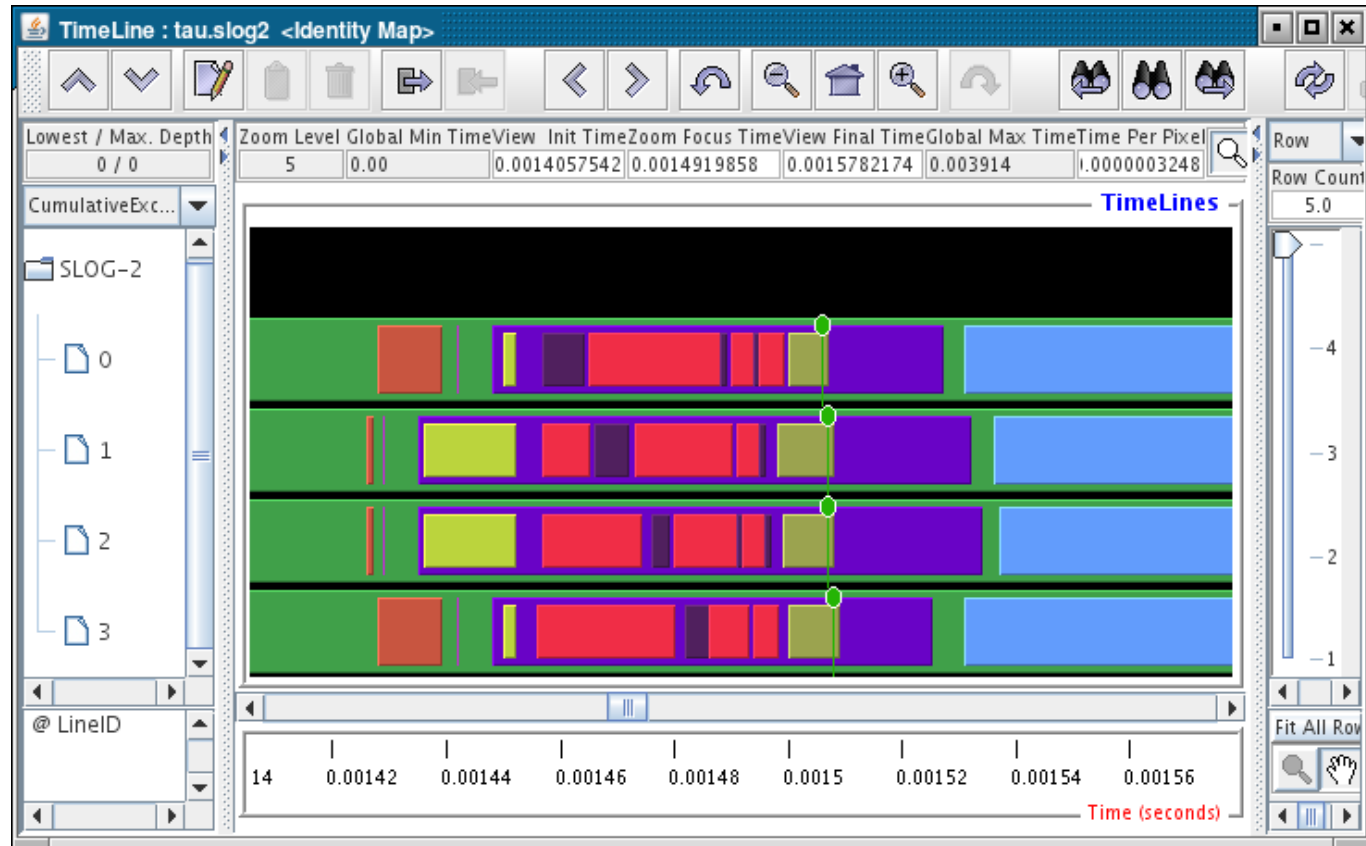


Visualization of the profile data with ParaProf.

Source: <http://www.cs.uoregon.edu/research/tau/docs/newguide/bk01ch04s02.html>



# TAU: Jumpshot



Visualization of trace data (trace of events in **time-line**) with Jumpshot  
Source: <http://www.cs.uoregon.edu/research/tau/docs/newguide/bk01ch04s03.html>

