

2DV605

Parallel Computing

GPU Programming

Suejb Memeti and Sabri Pllana

Department of Computer Science

Building D, Room D2236C, LNU, Växjö

<https://sabripllana.eu/>



Outline

- ❑ Introduction
- ❑ CUDA
- ❑ Example: Matrix Multiplication



Computing Platforms

- ❑ **Multi-core – seeks to maintain the speed of sequential programs while moving into multiple cores**
 - Consists of few cores optimized for sequential processing
 - Known as - **Host**
 - ex: Intel Multi-core CPUs
- ❑ **Many-core – focuses on the execution throughput of parallel applications**
 - Consists of thousands of smaller cores optimized for parallel execution
 - Peak floating point performance is about 10x faster compared to the performance of multi-core CPUs.
 - Known as - **Device**
 - ex: NVIDIA GPU, Intel Xeon Phi, FPGA
- ❑ **This performance gap has motivated developers to move the computationally intensive parts of their applications to the GPUs**



Credit: Intel

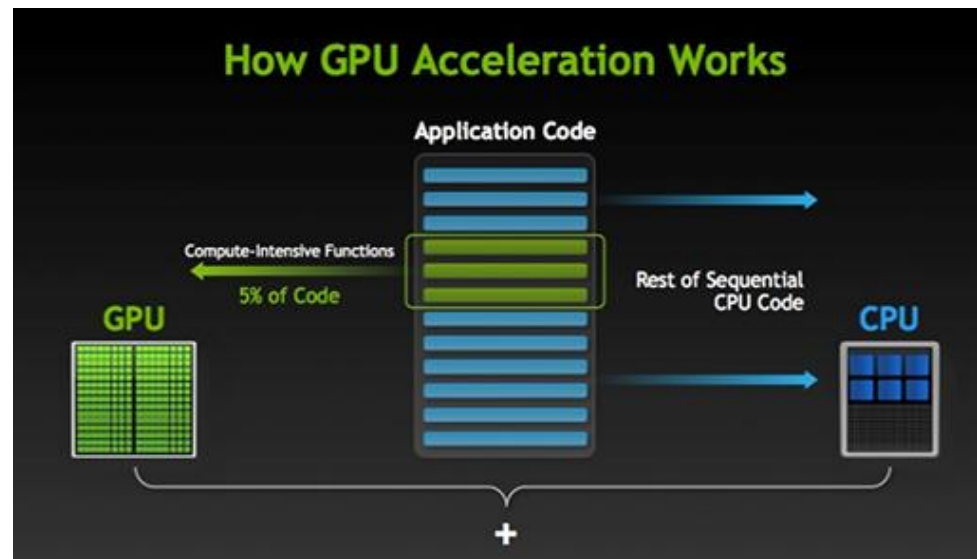


Credit: NVIDIA



Heterogeneous / Accelerated Computing

- ❑ Parts of the code (usually compute intensive tasks) are executed on the GPU
 - Known as **Offloading**
- ❑ The remainder of the code still runs on the CPU

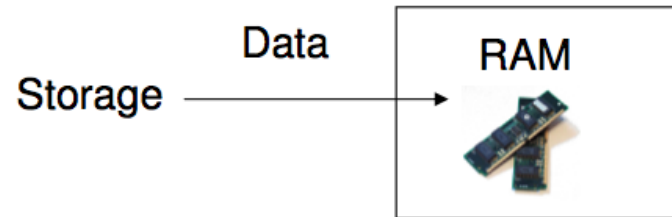


Credit: NVIDIA

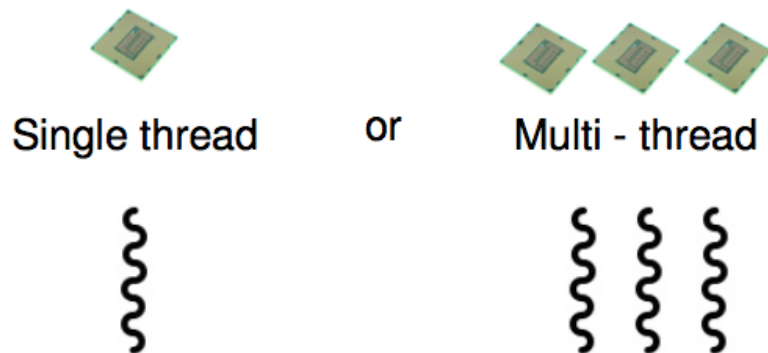


Homogeneous Computing

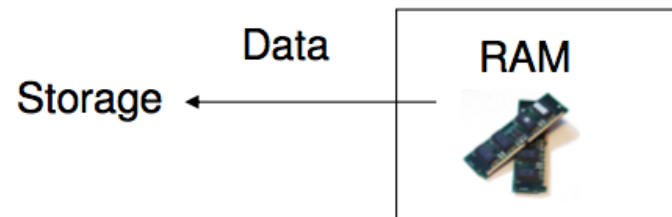
- ❑ Load data to the RAM



- ❑ Process data with CPU computation cores

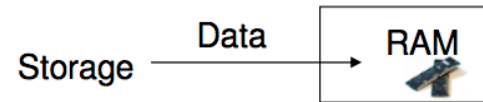


- ❑ Store result



Heterogeneous Computing

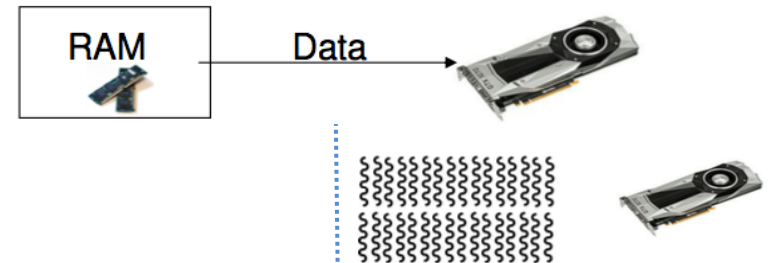
- ❑ Load data to the RAM



- ❑ Pre-process on CPU



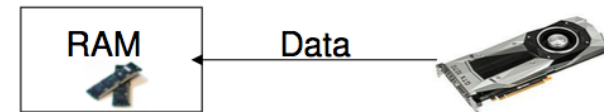
- ❑ Load data to the GPU memory



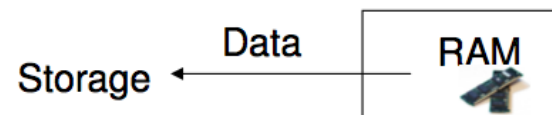
- ❑ Process data with GPU
 - CPU can either wait or execute some job



- ❑ Copy data back from GPU to RAM

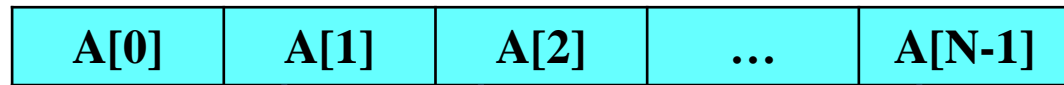


- ❑ Store result

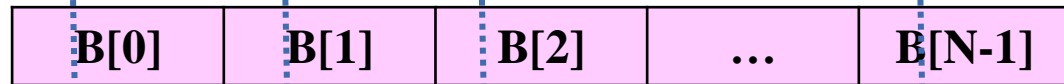


Data Parallelism – Vector Addition Example

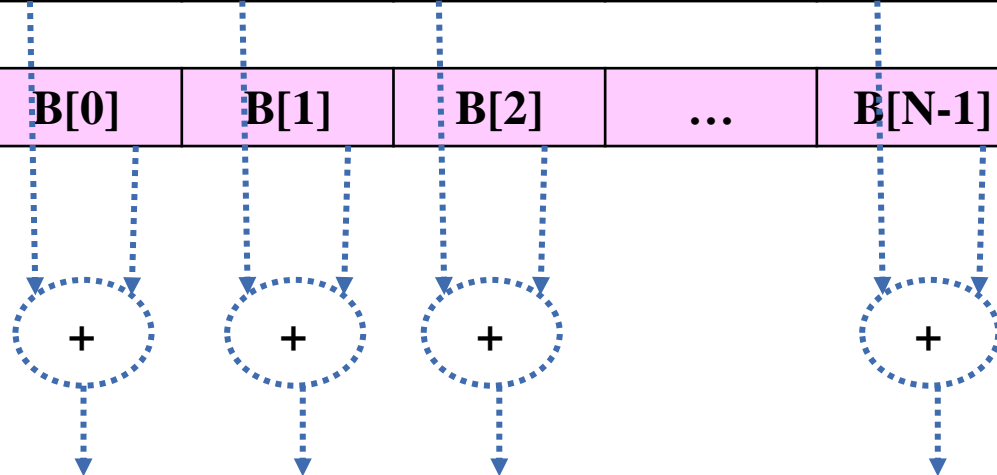
❑ Vector A



❑ Vector B

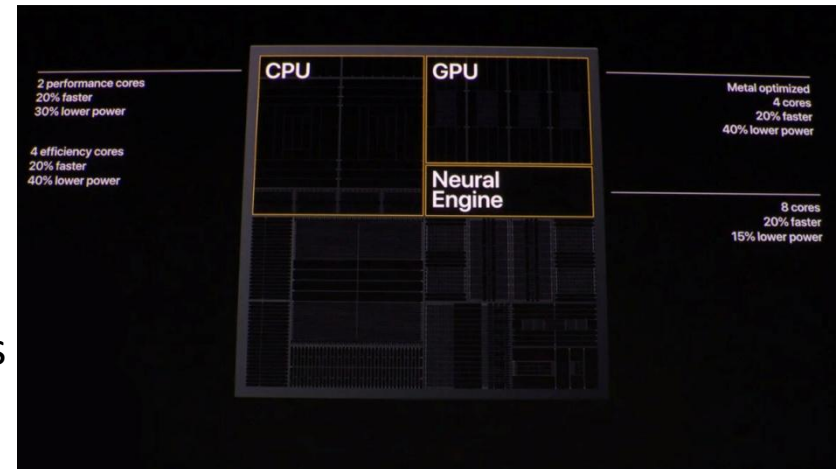


❑ Vector C



Graphics Processing Unit (GPU)

- ❑ **GPUs are everywhere**
 - smart-phones
 - supercomputers
- ❑ **Traditionally used for graphics rendering**
 - Computer games
 - Other 3D rendering applications
- ❑ **General Purpose GPUs**
 - Can be used to accelerate various scientific and engineering applications

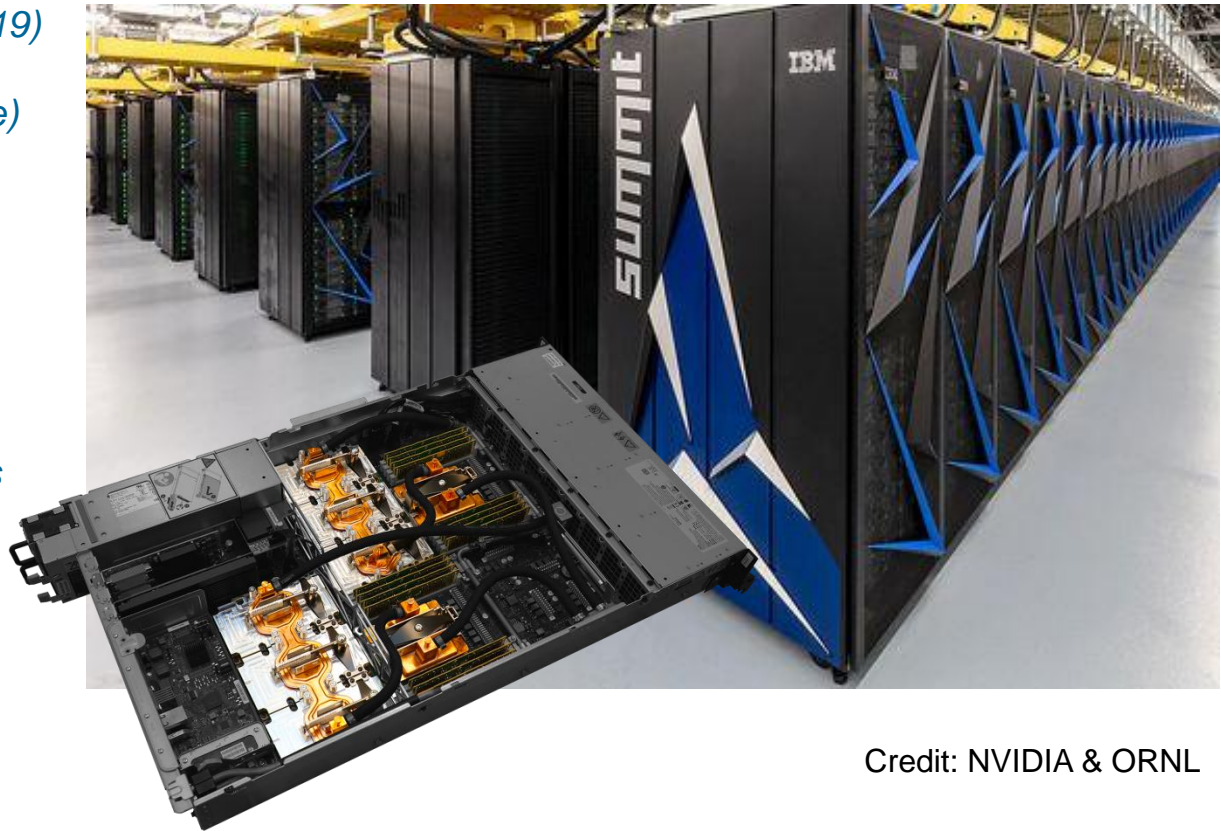


Apple A13 Bionic processor is used in iPhone 11



Supercomputer Example: Summit at ORNL

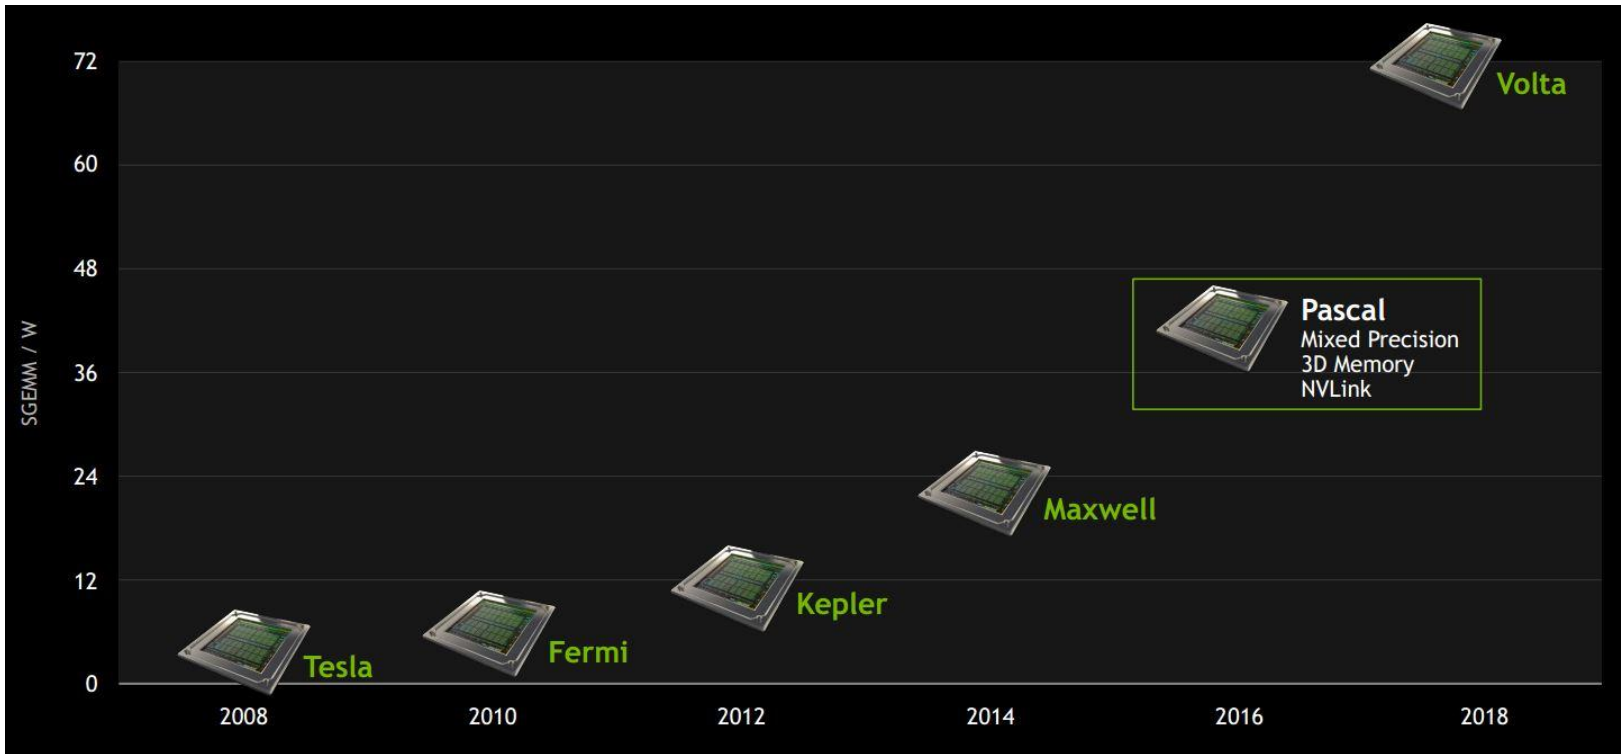
- *Rank 1 in TOP500 list (June 2019)*
- *200 petaflops (max performance)*
- *250 petabytes of data storage*
 - *74 years of HD video*
- *4608 computing nodes*
- *each node comprises*
 - *two IBM Power9 22-core CPUs*
 - *six NVIDIA Volta GPUs*
- *340 tons*
- *space of two tennis courts*



Credit: NVIDIA & ORNL



GPU Roadmap



Single precision floating General Matrix Multiply (SGEMM), 16K x 16K SGEMM

Credit: NVIDIA



Ida @ LNU

❑ Host

- 2 x Intel Xeon E5-2650 v4 CPUs
- 2.2– 2.9 GHz
- **12 Cores per CPU**
- 24 Threads per CPU
- 30MB cache
- Max. Memory Bandwidth is 76.8GB/s
- 384GB Memory
- 105W TDP

❑ Device

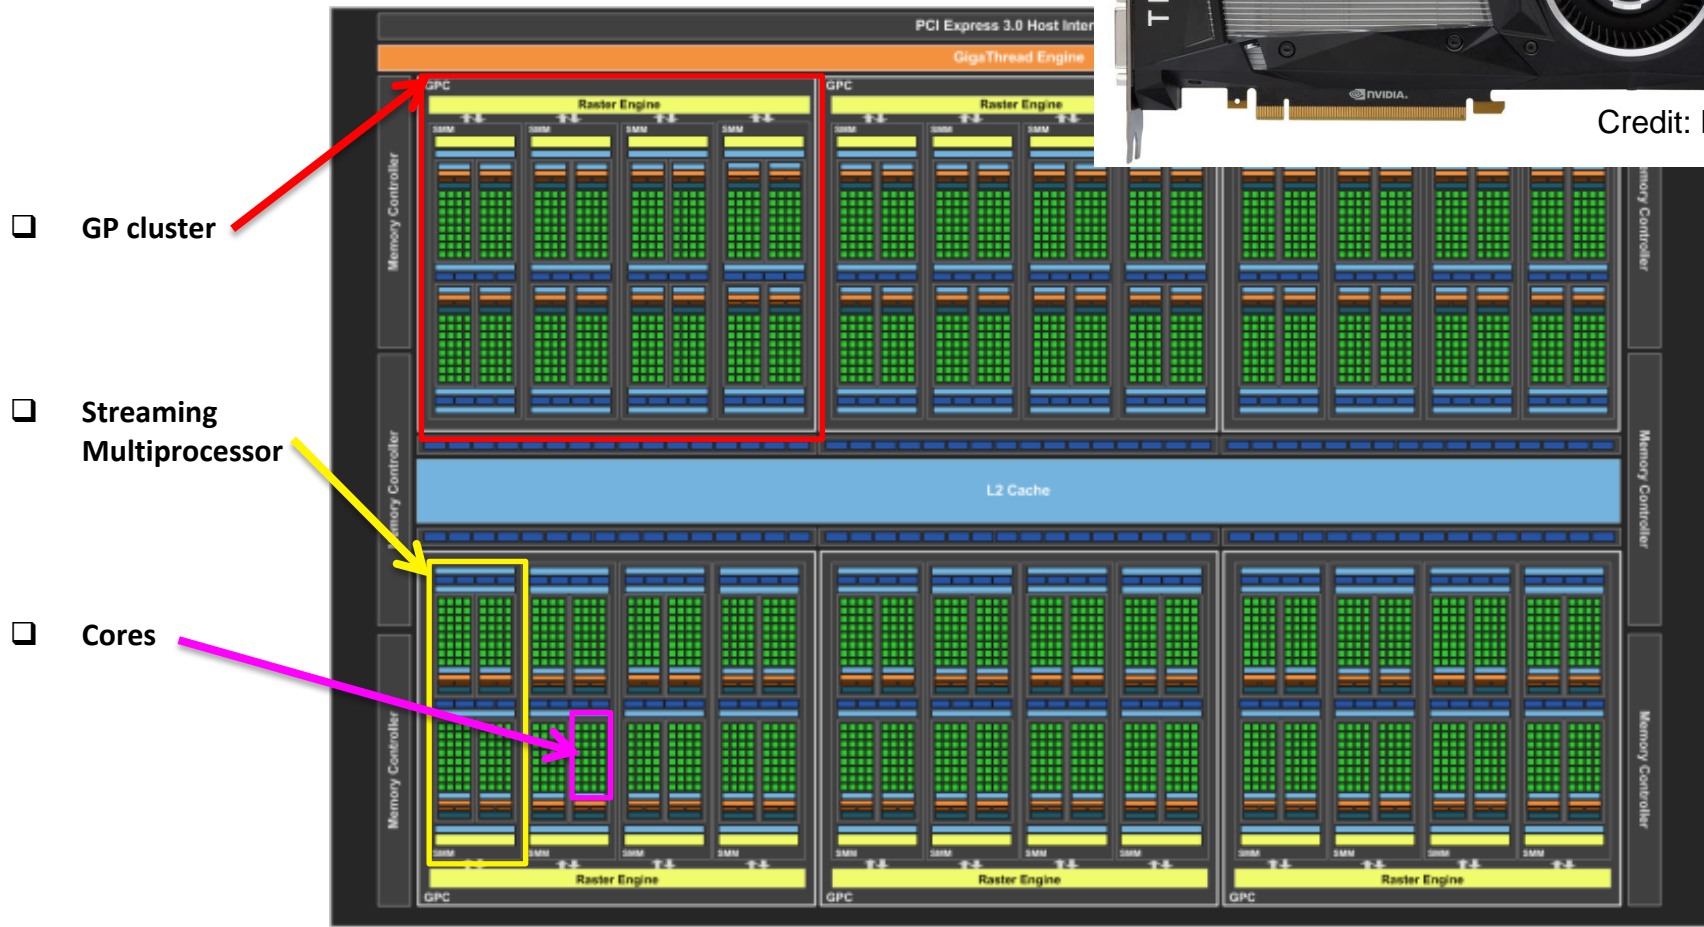
- 1 x GeForce GTX Titan X GPU
- 1 – 1.1 GHz
- 24 SMs
- **3072 cores**
- 336.5 GB/s Memory Bandwidth
- 12 GB Memory
- 250W TDP
- Warp Size is 32
- Maximum threads per block 1024



GTX Titan X



Credit: NVIDIA



CUDA



CUDA

- ❑ **A heterogeneous parallel programming interface that enables exploitation of data parallelism**
 - Hierarchical thread organization
 - Main interfaces for launching parallel execution
 - Mapping of thread index to data index

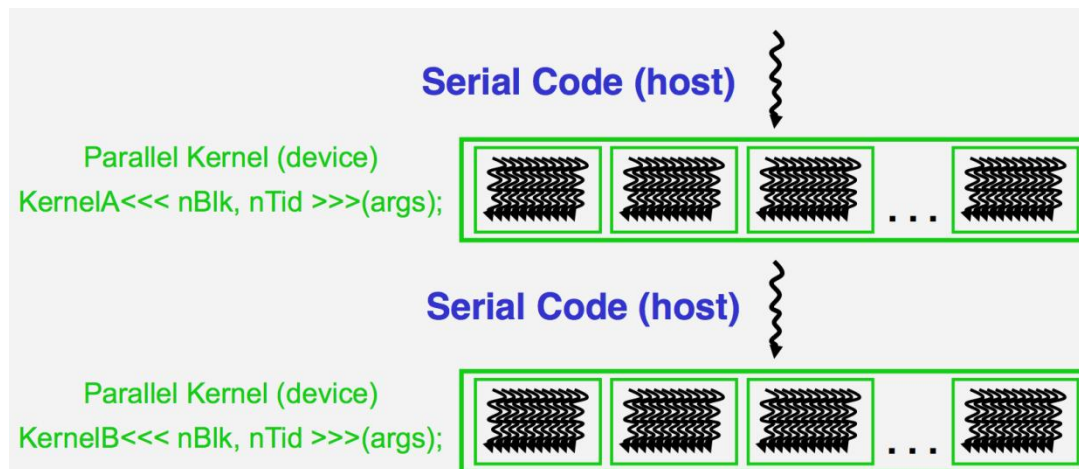
- ❑ **CUDA version installed on Ida @ LNU**
 - 8.0



Execution Model

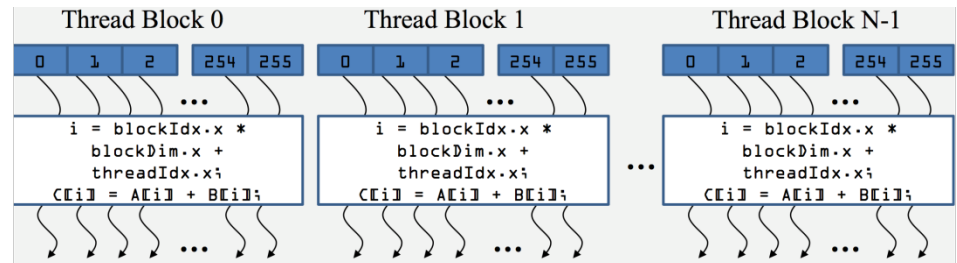
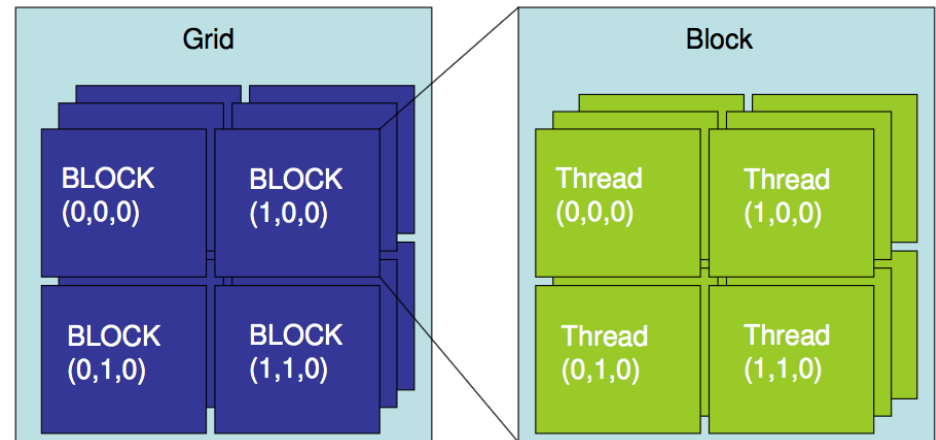
❑ Heterogeneous (host + device) C application

- Serial parts of the code are executed on the host
- Parallel parts of the code (CUDA kernels) are executed on the device



Grid, Block, Thread

- ❑ Grid – 3D Matrix of blocks
- ❑ Block – 3D Matrix of threads
- ❑ Thread – computation unit
 - Threads within a block communicate via the shared memory, atomic operations, and barrier synchronization
 - Threads in different blocks do not interact

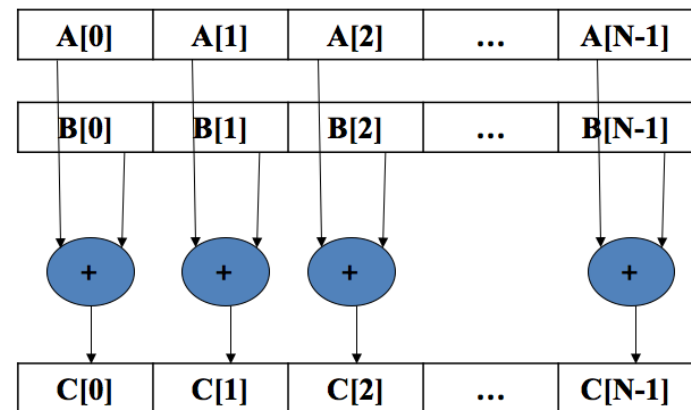


Computation Kernels

- ❑ A CUDA kernel is executed by an array of threads (so called grid)
 - All threads within a grid run the same kernel code
 - Each thread has indexes, which are used to compute the memory addresses that enable to access the data items

- ❑ The vector addition example

- We determine the index using the built-in variables:
 - `gridDim`, `blockIdx`, `blockDim`, `threadIdx`, `warpSize`
- Ex:
 - $i = blockDim.x * blockIdx.x + threadIdx.x;$
 $C[i] = A[i] + B[i];$



Grid, Block, Thread

- ❑ Any call to a CUDA kernel must specify the execution configuration
 - grid size
 - block size
 - number of threads per block `blockDim`
 - number of blocks per grid `gridDim`

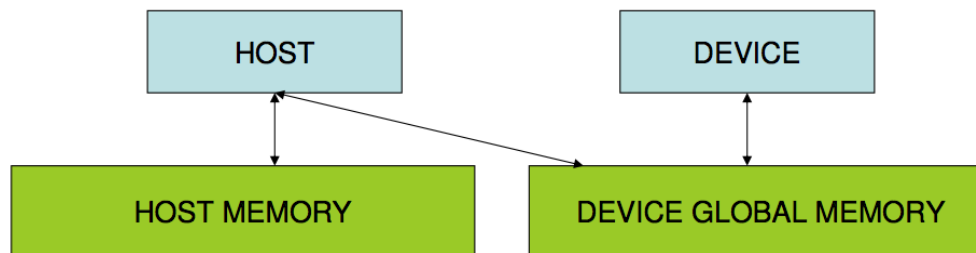
```
__global__ void vecAdd(double *a, double *b, double *c, int n) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    if(i < n) // make sure we do not go out of bounds  
        c[i] = a[i] + b[i];  
}  
// executing the kernel  
blockSize = 1024;  
gridSize = (int)ceil((float)n/blockSize);  
VecAdd<<<gridSize, blockSize>>>(a, b, c, n);
```



Memory Model

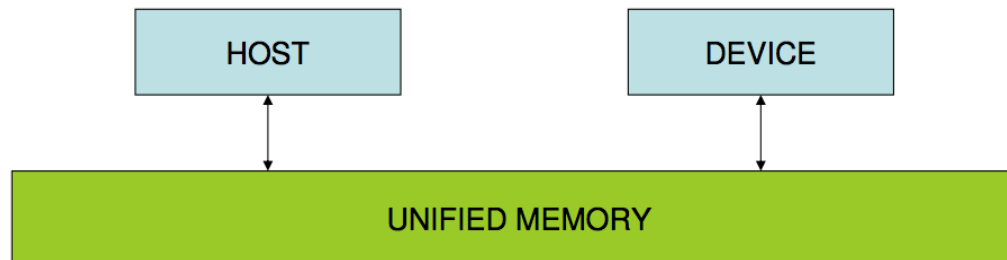
❑ NUMA

- Non-unified memory model



❑ Unified Memory Model

- From CUDA 6.0



Memory Management

❑ **cudaMalloc** – allocates memory on the device

❑ **cudaMemcpy** – copies data between host and device

▪ cudaMemcpyKind – CUDA memory copy types

- cudaMemcpyHostToHost = 0 Host → Host
- cudaMemcpyHostToDevice = 1 Host → Device
- cudaMemcpyDeviceToHost = 2 Device → Host
- cudaMemcpyDeviceToDevice = 3 Device → Device
- cudaMemcpyDefault = 4 Default based unified virtual address space

❑ **cudaFree** – frees the memory on the device



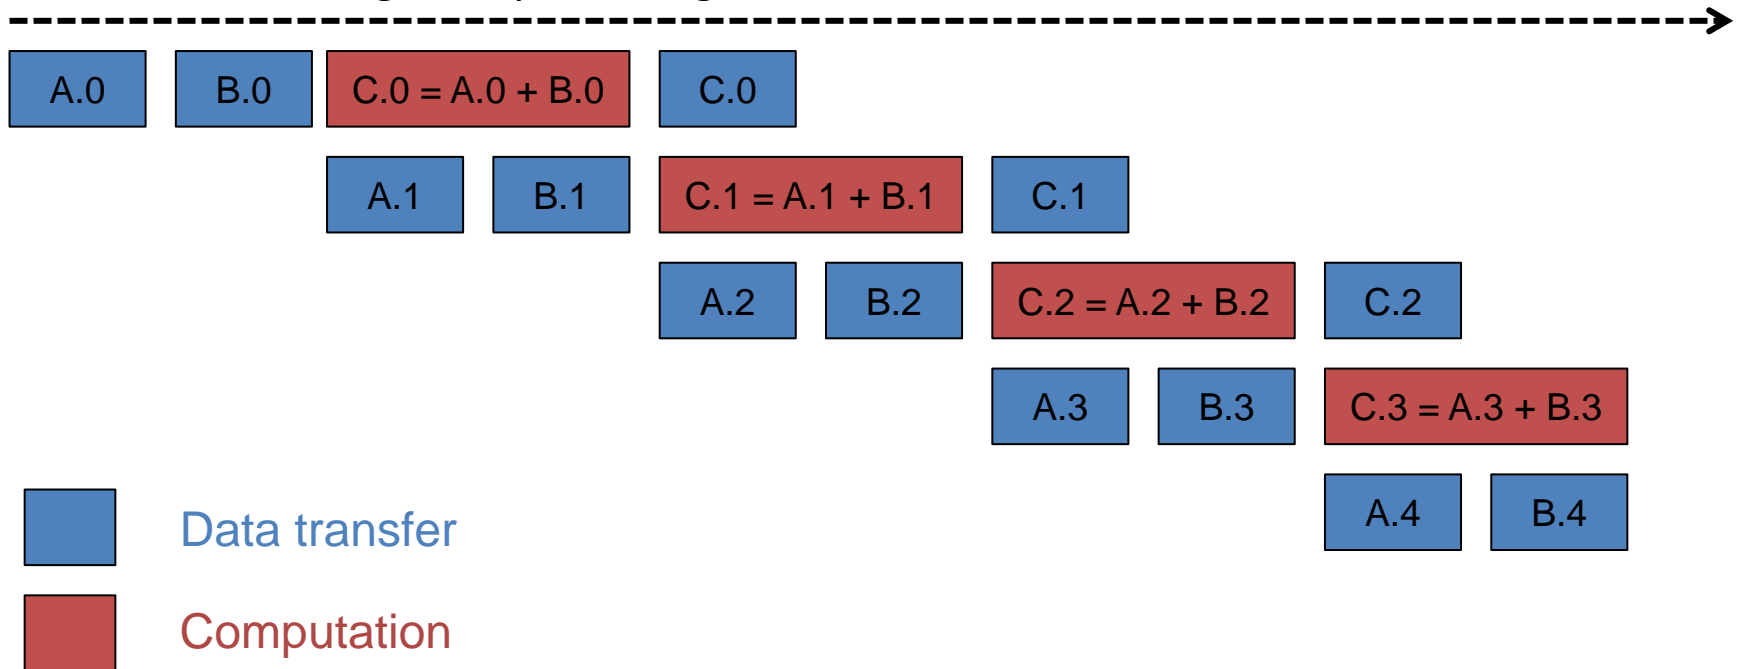
Basic Memory Management Example

```
...  
float *x, *y, *device_x, *device_y; // declare variables  
...  
x = (float*)malloc(N*sizeof(float)); // allocate memory on host for x  
y = (float*)malloc(N*sizeof(float)); // allocate memory on host for y  
...  
cudaMalloc(&device_x, N*sizeof(float)); // allocate memory on device for x  
cudaMalloc(&device_y, N*sizeof(float)); // allocate memory on device for y  
...  
cudaMemcpy(device_x, x, N*sizeof(float), cudaMemcpyHostToDevice); // copy x from  
host to device  
cudaMemcpy(device_y, y, N*sizeof(float), cudaMemcpyHostToDevice); // copy y from  
host to device  
... // run the computation and compute the result  
...  
cudaMemcpy(y, device_y, N*sizeof(float), cudaMemcpyDeviceToHost); // copy y from  
device to host  
... // store or print the result  
cudaFree(device_x); // free the memory for x on device  
cudaFree(device_y); // free the memory for y on device
```



Asynchronous Transfer

- ❑ Overlap computation and communication
- ❑ Data domain decomposition
 - Divide large arrays into segments



Example: Matrix Multiplication



Environment Variables

- ❑ NVIDIA CUDA Compiler (NVCC) 8.0 is installed in

[/usr/local/cuda-8.0/bin/nvcc/](#)

- Set the following environment variables

- PATH:

`export PATH=$PATH:/usr/local/cuda-8.0/bin:$PATH`

- LIBRARY_PATH:

`LIBRARY_PATH=/usr/local/cuda-8.0/lib:/usr/local/cuda-8.0/lib64:$LIBRARY_PATH`

- LD_LIBRARY_PATH:

`export LD_LIBRARY_PATH=/usr/local/lib:/usr/lib:/usr/local/lib64:/usr/lib64:/usr/local/cuda-8.0/lib:$LD_LIBRARY_PATH`

- You can put them in [/home/username/.profile](#) file
 - To automatically load on user login



Example of .profile File

```
export PATH=$PATH:/usr/local/cuda-8.0/bin:$PATH
```

```
export LIBRARY_PATH=/usr/local/cuda-8.0/lib:/usr/local/cuda-8.0/lib64/:$LIBRARY_PATH
```

```
export LD_LIBRARY_PATH=/usr/local/cuda-8.0/lib64/:$LD_LIBRARY_PATH
```



Initialization

```
#define TILE_SIZE 16
#define MATRIX_SIZE 512

float * hostA; // The A matrix
float * hostB; // The B matrix
float * hostC; // The output C matrix
float * deviceA;
float * deviceB;
float * deviceC;

hostA = (float*) malloc(sizeof(float) * n * n);
...
for(int i = 0; i < n; ++i) {
    hostA[i] = rand()/float(RAND_MAX)*24.f+1.f;
    hostB[i] = rand()/float(RAND_MAX)*24.f+1.f;
    hostC[i] = 0;
}
```



Data Transfer To / From Device

```
// allocate device memory
cudaError_t err1 = cudaMalloc((void **) &deviceA, sizeA);
cudaError_t err2 = cudaMalloc((void **) &deviceB, sizeB);
cudaError_t err3 = cudaMalloc((void **) &deviceC, sizeC);

// Copy memory to the GPU
cudaMemcpy(deviceA, hostA, sizeA, cudaMemcpyHostToDevice);
cudaMemcpy(deviceB, hostB, sizeB, cudaMemcpyHostToDevice);
cudaMemcpy(deviceC, hostC, sizeC, cudaMemcpyHostToDevice);

// execute the kernel

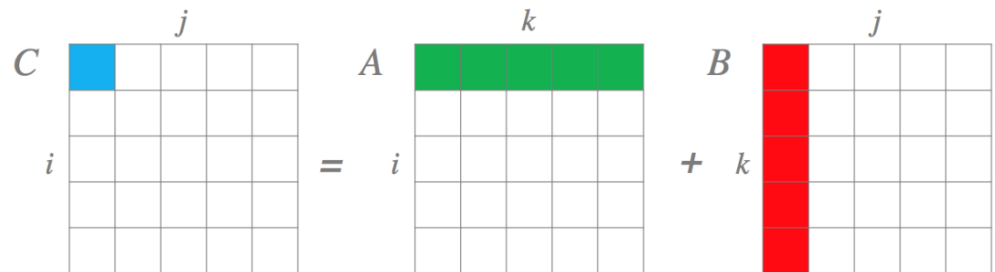
// Copy the GPU memory back to the CPU
cudaMemcpy(hostC, deviceC, sizeC, cudaMemcpyDeviceToHost);

// Free the GPU memory
cudaFree(deviceA);
cudaFree(deviceB);
cudaFree(deviceC);
```



Matrix Multiplication on Host

```
void mmultHost(float * A, float * B, float * C, int numRows,
int numACols, int numBRows, int numBCols, int numCRows, int
numCCols) {
    for (int i=0; i < numRows; i++) {
        for (int j = 0; j < numACols; j++) {
            for (int k = 0; k < numCCols; k++) {
                C[i*numCCols+j] += A[i*numACols+k] * B[k*numBCols+j];
            }
        }
    }
}
```



Matrix Multiplication on Device

```
dimBlock(TILE_SIZE, TILE_SIZE, 1);  
dim3 dimGrid((numCCs/TILE_SIZE) + 1, (numCRows/TILE_SIZE) + 1, 1);  
mmultDevice<<<dimGrid, dimBlock>>>(deviceA, deviceB, deviceC, numRows,  
numACols, numBRows, numBCols, numCRows, numCCols);
```

```
__global__ void mmultDevice(float *A, float *B, float *C, int numRows,  
int numACols, int numBRows, int numBCs, int numCRows, int numCCols) {  
    int row = blockIdx.y * blockDim.y + threadIdx.y;  
    int col = blockIdx.x * blockDim.x + threadIdx.x;  
    if(numACols != numBRows) { return; }  
    if((row < numRows) && (col < numBCols)) {  
        float Cvalue = 0;  
        for(int i = 0; i < numACs; ++i) {  
            Cvalue += A[row * numACols + i] * B[i * numBCols + col];  
        }  
        C[row * numCColumns + col] = Cvalue;  
    }  
}
```



Compile, Execute, Monitor

☐ Compile

- `nvcc file_name.cu -o file_name.o`

☐ Execute

- `./file_name.o {parameters}`

☐ Monitor

- `nvidia-smi`
- `Watch -n .1 nvidia-smi`



Literature

- ❑ Programming Massively Parallel Processors, Third Edition: A Hands-on Approach, 2017 by David B. Kirk and Wen-mei W. Hwu
- ❑ Heterogeneous System Architecture: A New Compute Platform Infrastructure 1st Edition by Wen-mei W. Hwu
- ❑ Principles of Parallel Programming by Calvin Lin and Lawrence Snyder, Pearson Addison Wesley
- ❑ CUDA Application Design and Development by Rob Farber
- ❑ CUDA by Example: An Introduction to General-Purpose GPU Programming, by Jason Sanders and Edward Kandrot.
- ❑ Patterns for Parallel Programming by Timothy G. Mattson, Beverly A. Sanders, and Berna L. Massingill, Addison Wesley
- ❑ CUDA Toolkit reference manual <http://docs.nvidia.com/cuda>
- ❑ CUDA C Programming guide <http://docs.nvidia.com/cuda/cuda-c-programming-guide>

