

2DV605

Parallel Computing

OpenMP

Sabri Pllana, PhD

Department of Computer Science

Building D, Room D2236C, LNU, Växjö

sabri.pllana@lnu.se, <http://homepage.lnu.se/staff/saplaa/>



Outline

- ☐ Introduction
- ☐ Example: Hello World
- ☐ Compiler Directives
- ☐ Runtime Library Routines
- ☐ Environment Variables
- ☐ Example: Matrix Multiplication



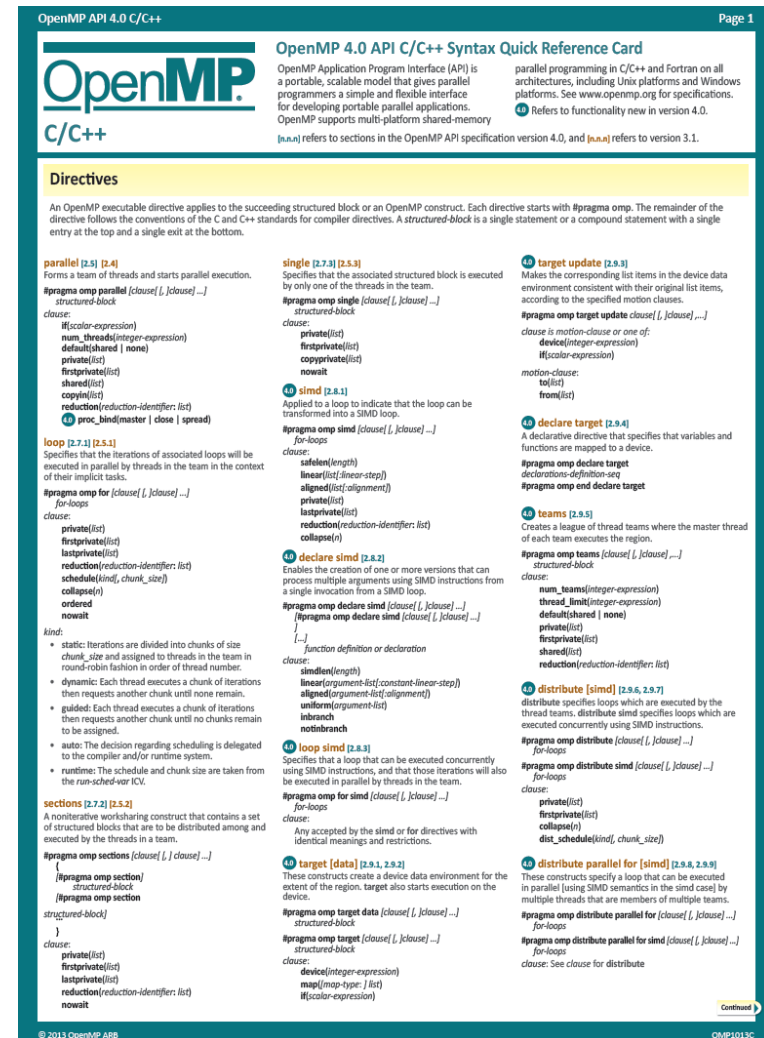
Selected Literature

❑ OpenMP Specifications

- Version 4.0, July 2013
(supported by GCC 4.9 or later version)
- OpenMP 4.5, November 2015
- <http://openmp.org/wp/openmp-specifications/>

❑ OpenMP 4.0 API C/C++ Syntax Quick Reference Card

- <http://openmp.org/mp-documents/OpenMP-4.0-C.pdf>



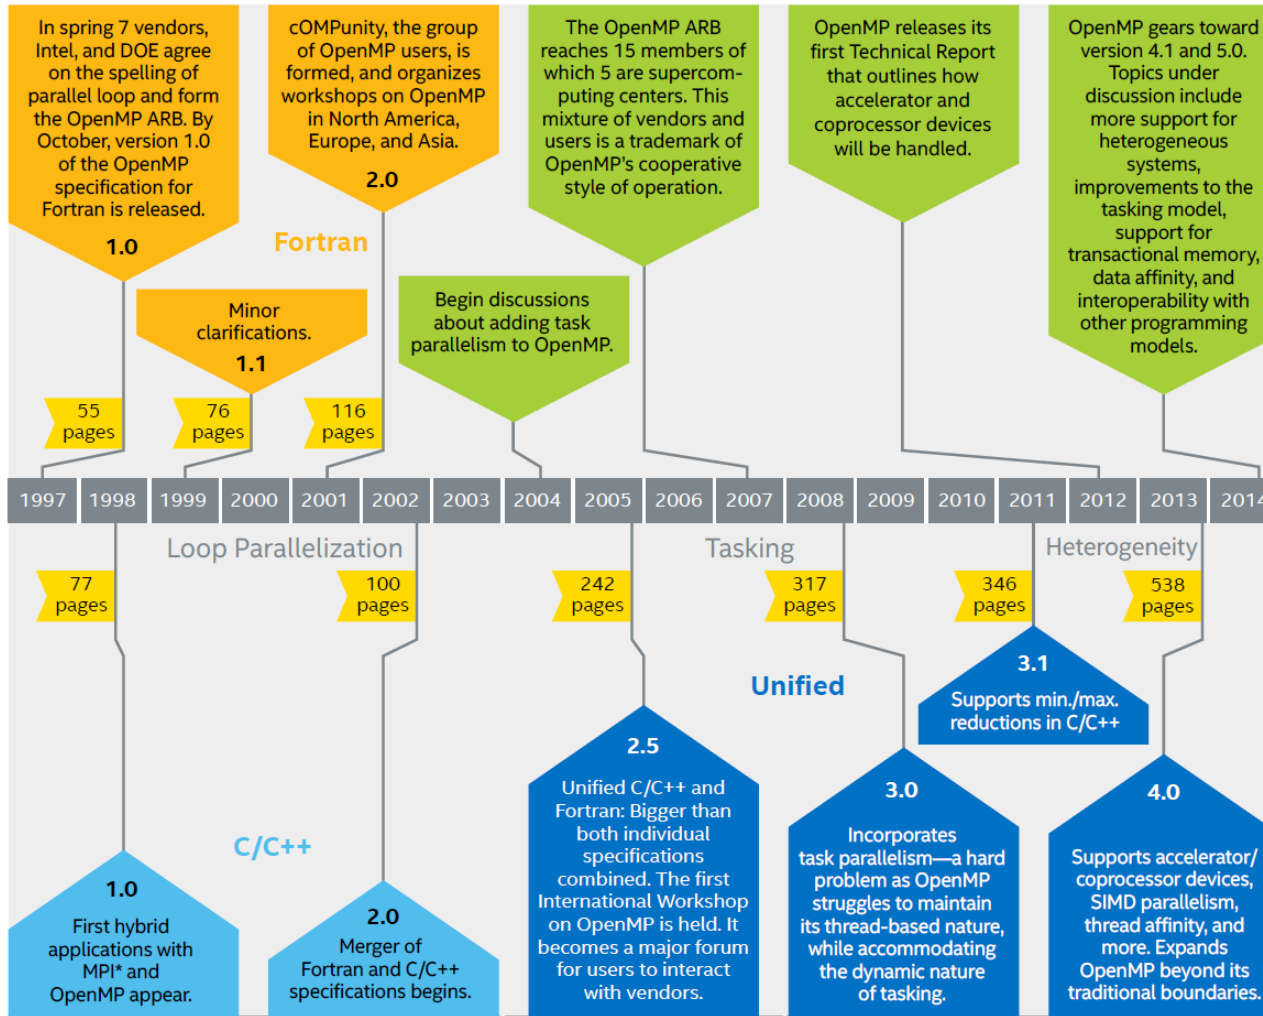
Introduction



OpenMP History

1996 Vendors provide similar but different solutions for loop [parallelism](#), causing portability and maintenance problems.

Kuck and Associates, Inc. (KAI) | SGI | Cray | IBM | High Performance Fortran (HPF) | Parallel Computing Forum (PCF)



Credit: The Parallel Universe, June 2014, © Intel Corporation



OpenMP Architecture Review Board (ARB)

ARB oversees OpenMP specification

*Permanent ARB members
develop and sell OpenMP products*

- AMD (Greg Stoner)
- ARM (Chris Adeniyi-Jones)
- Cray (Luiz DeRose)
- Fujitsu (Eiji Yamanaka)
- HP (Sujoy Saraswati)
- IBM (Kelvin Li)
- Intel (Xinmin Tian)
- Micron (Kirby Collins)
- NEC (Kazuhiro Kusano)
- NVIDIA (Jeff Larkin)
- Oracle Corporation (Nawal Copt)
- Red Hat (Matt Newsome)
- Texas Instruments (Eric Stotzer)

Auxiliary ARB members do not sell OpenMP products

- Argonne National Laboratory (Kalyan Kumaran)
- ASC/Lawrence Livermore National Laboratory (B. R. de Supinski)
- Barcelona Supercomputing Center (Xavier Martorell)
- Bristol University (Simon McIntosh-Smith)
- cOMPunity (Barbara Chapman/Yonghong Yan)
- Edinburgh Parallel Computing Centre (EPCC) (Mark Bull)
- INRIA (Olivier Aumage)
- Los Alamos National Laboratory (David Montoya)
- Lawrence Berkeley National Laboratory (Alice Koniges/Helen He)
- NASA (Henry Jin)
- Oak Ridge National Laboratory (Oscar Hernandez)
- RWTH Aachen University (Dieter an Mey)
- Sandia National Laboratory (Stephen Olivier)
- Texas Advanced Computing Center (Kent Milfeld)
- University of Houston (Deepak Eachempati/Jeremy Kemp)

Status: September 2016



OpenMP Overview

- ❑ **OpenMP supports shared-memory parallelism**
 - extends C, C++, Fortran programming languages

- ❑ **OpenMP provides constructs for**
 - single program multiple data (SPMD)
 - tasks
 - device
 - worksharing and synchronization


- ❑ **OpenMP (Application Programming Interface) API**
 - compiler directives
 - library routines
 - environment variables



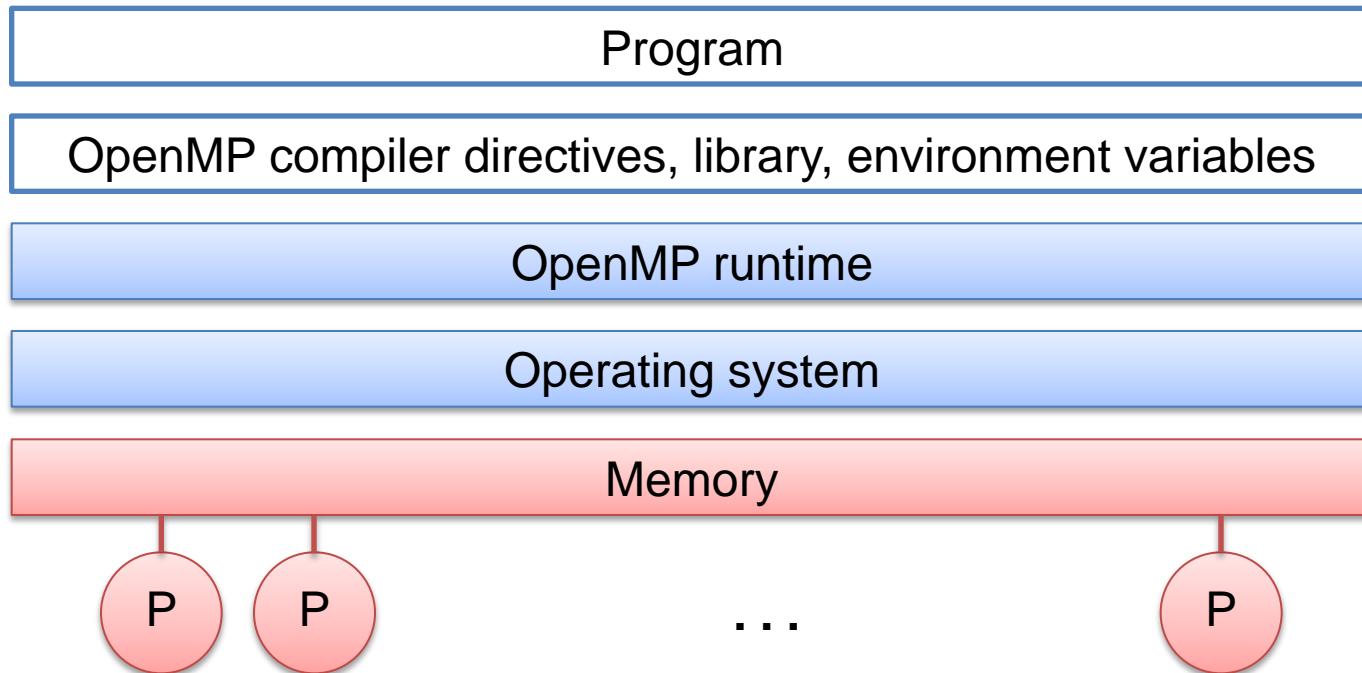
Threads

- ❑ **Edward A. Lee (2006): The Problem with Threads**
 - *“Threads, as a model of computation, are wildly nondeterministic, and the job of the programmer becomes one of pruning that **nondeterminism**”*
 - *“non-trivial multi-threaded programs are **incomprehensible to humans**”*
 - *“Threads must be relegated to the engine room of computing, to be suffered only by expert technology providers”*

non determinism: various outputs may result for the same input



OpenMP System Stack



OpenMP and Threads

```
#pragma omp parallel num_threads(4)
{
    foobar ();
}
```



```
void thunk ()
{
    foobar ();
}

pthread_t tid[4];
for (int i = 1; i < 4; ++i)
    pthread_create (
        &tid[i], 0, thunk, 0);
thunk();

for (int i = 1; i < 4; ++i)
    pthread_join (tid[i]);
```

Credit: Tim Mattson




OpenMP Execution Model

❑ Fork and join model

- parallel regions may be nested
- implicit barrier at the end of parallel construct
- a program may comprise an arbitrary number of parallel constructs

❑ No guarantee for synchronous file I/O during parallel execution

- programmer is responsible for synchronization, if multiple threads access the same file



Example: Hello World



Hello World: C-Code Extended with OpenMP

```
#include <omp.h>
#include <stdio.h>
main () {
    int threads, id;
    #pragma omp parallel private(id)
    { // Begin of parallel region
        id = omp_get_thread_num(); // Get the thread ID
        printf("Hello World, I am thread = %d\n", id);
        // The thread with ID = 0 outputs the total number of threads
        if (id == 0) {
            threads = omp_get_num_threads();
            printf("Number of threads = %d\n", threads);
        }
    } // End of parallel region
}
```



Hello World: Compilation and Execution

Compilation

```
-bash-4.1$ gcc -o hello -fopenmp hello.c
```

Execution on “Emil”

```
-bash-4.1$ ./hello
```

*“Emil” comprises two processors of the type
Intel Xeon E5-2695v2,
each has 12 cores, and supports two threads per core
(two logical cores per physical core).*

In total are 48 threads.

```
Hello World, I am thread = 10
Hello World, I am thread = 12
...
Hello World, I am thread = 29
Hello World, I am thread = 15
Hello World, I am thread = 33
Hello World, I am thread = 28
Hello World, I am thread = 42
Hello World, I am thread = 8
Hello World, I am thread = 17
Hello World, I am thread = 18
Hello World, I am thread = 14
Hello World, I am thread = 47
Hello World, I am thread = 41
Hello World, I am thread = 44
Hello World, I am thread = 38
Hello World, I am thread = 37
Hello World, I am thread = 0
Number of threads = 48
Hello World, I am thread = 16
Hello World, I am thread = 22
Hello World, I am thread = 2
Hello World, I am thread = 1
Hello World, I am thread = 6
Hello World, I am thread = 26
```



Emil: CPU Information

```
-bash-4.1$ cat /proc/cpuinfo | more
processor          : 0
vendor_id         : GenuineIntel
cpu family        : 6
model             : 62
model name        : Intel(R) Xeon(R) CPU E5-2695 v2 @ 2.40GHz
stepping          : 4
cpu MHz           : 2399.828
cache size        : 30720 KB
...
processor          : 47
vendor_id         : GenuineIntel
cpu family        : 6
model             : 62
model name        : Intel(R) Xeon(R) CPU E5-2695 v2 @ 2.40GHz
stepping          : 4
cpu MHz           : 2399.828
cache size        : 30720 KB
...
```



Environment Variable OMP_NUM_THREADS

Set the number of OpenMP threads to 12

```
-bash-4.1$ export OMP_NUM_THREADS=12
```

```
-bash-4.1$ ./hello
Hello World, I am thread = 10
Hello World, I am thread = 5
Hello World, I am thread = 2
Hello World, I am thread = 7
Hello World, I am thread = 6
Hello World, I am thread = 3
Hello World, I am thread = 4
Hello World, I am thread = 0
Number of threads = 12
Hello World, I am thread = 9
Hello World, I am thread = 11
Hello World, I am thread = 1
Hello World, I am thread = 8
```

```
-bash-4.1$ ./hello
Hello World, I am thread = 6
Hello World, I am thread = 11
Hello World, I am thread = 2
Hello World, I am thread = 4
Hello World, I am thread = 7
Hello World, I am thread = 5
Hello World, I am thread = 9
Hello World, I am thread = 10
Hello World, I am thread = 3
Hello World, I am thread = 8
Hello World, I am thread = 1
Hello World, I am thread = 0
Number of threads = 12
```

```
-bash-4.1$ ./hello
Hello World, I am thread = 0
Number of threads = 12
Hello World, I am thread = 8
Hello World, I am thread = 3
Hello World, I am thread = 4
Hello World, I am thread = 6
Hello World, I am thread = 2
Hello World, I am thread = 9
Hello World, I am thread = 7
Hello World, I am thread = 1
Hello World, I am thread = 5
Hello World, I am thread = 10
Hello World, I am thread = 11
```

*Executed the same program three times;
the order of outputs varies.*



OpenMP Compiler Directives

Comprehensive information is available in OpenMP Specification



OpenMP Directive Format: C/C++

❑ OpenMP directives

- start with `#pragma omp`
- are case sensitive
- apply to the succeeding statement (structured block)

```
#pragma omp directive-name [clause[ [,] clause]...] new-line
```



Parallel Construct

❑ Creates a thread team and initiates parallel execution

❑ *clause* may be

- if(scalar-expression)
- `num_threads(integer-expression)` // determines the number of threads
- default(shared | none)
- private(list)
- firstprivate(list)
- shared(list)
- copyin(list)
- reduction(reduction-identifier :list)
- proc_bind(master | close | spread)

```
#pragma omp parallel [clause[ [, ]clause] ...] new-line  
structured-block
```



Worksharing Constructs

❑ Share the work of executing a region

- each thread of the team executes a part
- a worksharing construct is used in the context of `parallel` construct

❑ Worksharing constructs

- `loop`
- `sections`
- `single`

❑ Implicit barrier at the end of the worksharing construct

- `nowait` clause indicates that the implicit barrier may be omitted



Worksharing Constructs: Loop

- ❑ Iterations of the loop are executed in parallel
- ❑ *clause* may be
 - `private(list)`
 - `firstprivate(list)`
 - `lastprivate(list)`
 - `reduction(reduction-identifier: list)`
 - `schedule(kind[, chunk_size])` // default `chunk_size` is 1
 - `collapse(n)`
 - `ordered`
 - `nowait`

```
#pragma omp for [clause[[,] clause] ... ] new-line  
for-loops
```



Worksharing Constructs: Loop Schedule

- ❑ **schedule(*static*, *chunk_size*)**
 - divides iterations into chunks of size *chunk_size*
 - assigns chunks to the threads in the team
 - round-robin fashion
- ❑ **schedule(*dynamic*, *chunk_size*)**
 - schedules iterations based on thread requests
 - after the execution of a chunk the thread request the next chunk
- ❑ **schedule(*guided*, *chunk_size*)**
 - similar to dynamic; but, the chunk size decreases during scheduling
- ❑ **schedule(auto)**
 - delegates scheduling decision to compiler or/and run-time system
- ❑ **schedule(runtime)**
 - defers the scheduling decision to run-time



Worksharing Constructs: Sections

- ❑ A collection of structured blocks is executed in parallel
- ❑ *clause* may be
 - `private(list)`, `firstprivate(list)`, `lastprivate(list)`,
 - `reduction(reduction-identifier:list)`,
 - `nowait`

```
#pragma omp sections [clause[[,] clause] ...] new-line
{
    [#pragma omp section new-line]
    structured-block
    [#pragma omp section new-line]
    structured-block]
    ...
}
```



Worksharing Constructs: Single

- ❑ Indicates that the structured block is executed by only one thread
- ❑ *clause* may be
 - `private(list)`
 - `firstprivate(list)`
 - `copyprivate(list)`
 - `nowait`

```
#pragma omp single [clause[,] clause] ...] new-line  
structured-block
```



SIMD Construct

- ❑ **Specifies that the loop may be transformed into a SIMD loop**
 - execute loop iterations using SIMD instructions
 - SIMD instruction is a machine instruction that may operate on many data elements, SIMD stands for Single Instruction Multiple Data
 - *vector operation*: processes multiple pairs of operands in one step
- ❑ *clause* **may be**
 - safelen(length),
 - linear(list[:linear-step]), aligned(list[:alignment]),
 - private(list), lastprivate(list),
 - reduction(reduction-identifier:list), collapse(n)

```
#pragma omp simd [clause[,] clause] ...] new-line  
for-loops
```



Declare SIMD Construct

- ❑ A function processes multiple arguments using SIMD instructions

```
#pragma omp declare simd [clause[,] clause] ...] new-line  
[...]  
function definition or declaration
```

```
void findmin (int *a, int *b, int *c) {  
    #pragma omp simd  
    for (i=0; i<N; i++)  
        c[i] = min(a[i], b[i]);  
}  
  
#pragma omp declare simd  
int min (int a, int b) {  
    return a < b ? a : b;  
}
```

Example credit: IBM



Loop SIMD Construct

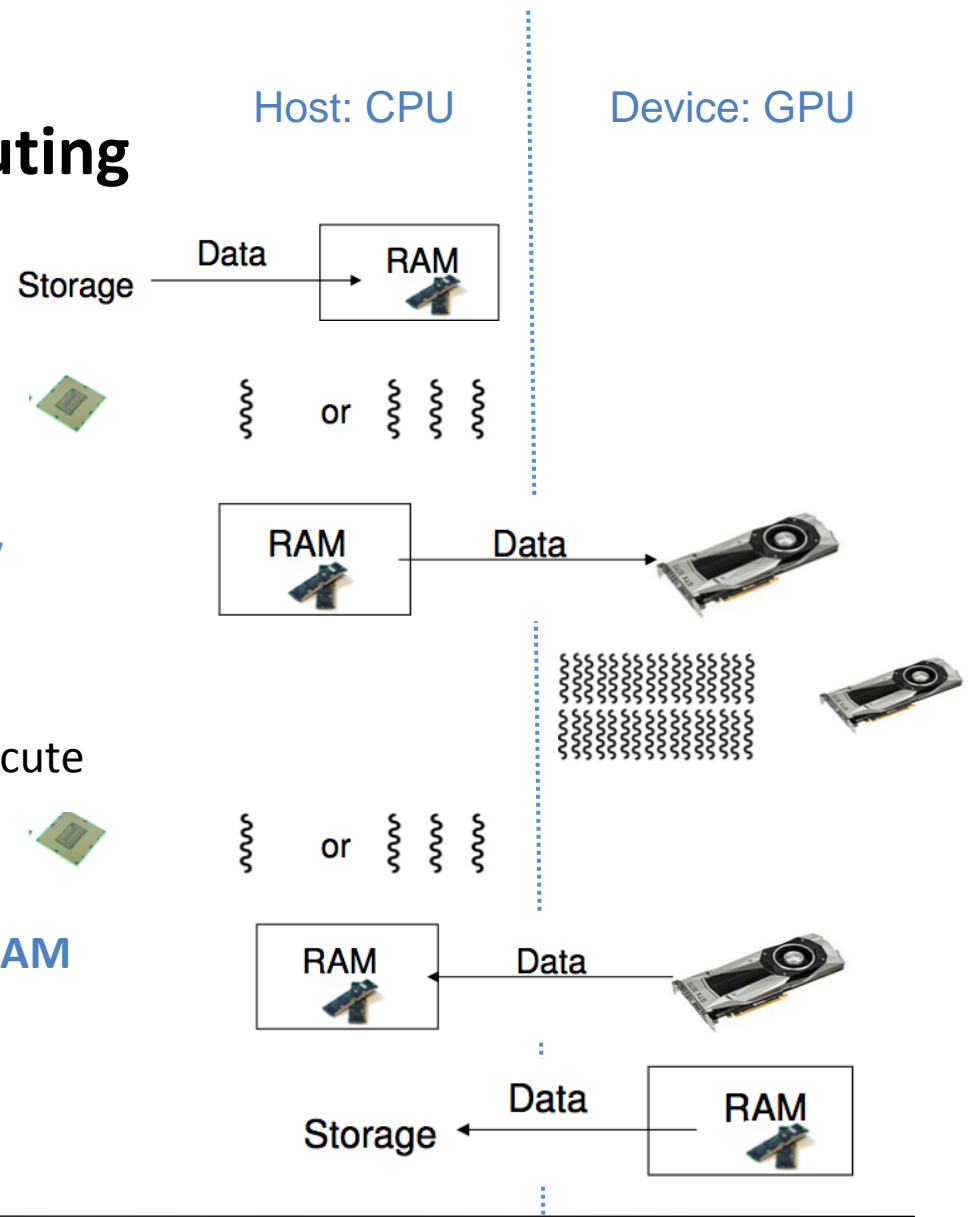
- ❑ Indicates a loop that may be executed using SIMD instructions
 - loop iterations are executed in parallel also by team threads
 - *first the loop iterations are distributed across the team of threads, and thereafter chunks of iterations are transformed into SIMD loop*
- ❑ *clause* **may be**
 - any **for** or **simd** clause

```
#pragma omp for simd [clause[,] clause] ...] new-line  
for-loops
```



Heterogeneous Computing

- ☐ Load data to the RAM
- ☐ Pre-process on CPU
- ☐ Load data to the GPU memory
- ☐ Process data with GPU
 - CPU can either wait or execute some job
- ☐ Copy data back from GPU to RAM
- ☐ Store result



Device Constructs: target data

❑ Creates the device data environment

- maps variables between *host* and *target device*
- examples of *target devices*: **GPU**, Intel Xeon Phi,..
- *host*: device that initiates OpenMP execution
- typically the *host* is **CPU**

❑ *clause* may be

- device (integer-expression)
- map ([map-type :] list) *// relates variables on host and target device*
- if (scalar-expression)

```
#pragma omp target data [clause[[,] clause],...] new-line  
structured-block
```



Device Constructs: target

- ❑ Creates the device data environment and initiates the execution on device

```
#pragma omp target [clause[,] clause],... new-line  
structured-block
```

```
#pragma omp target data  
  map(to:a,b,c)  
  map(from:e,f) {  
    #pragma omp target {  
      e = f(a,b,c);  
    }  
    #pragma omp target {  
      f = g(a,b,c);  
    }  
  }
```

Example credit: Intel



Device Constructs: target update

❑ Ensures consistency between data on host and target device

❑ *clause* may be

- *to(list)*
- *from(list)*
- *device(integer-expression)*
- *if(scalar-expression)*

```
#pragma omp target update [clause[[,] clause],...] new-line
```



Combined Constructs: Parallel Loop

- ❑ Combined constructs are shortcuts for specifying constructs
 - for instance: **parallel for**
- ❑ *clause* may be
 - any **parallel** or **for** clause

```
#pragma omp parallel for [clause[,] clause] ...] new-line  
for-loop
```



Tasking Constructs: task

❑ Defines a task

- generates a task from *structured-block*

❑ *clause* may be

- if(scalar-expression)
- final(scalar-expression)
- untied
- default(shared | none)
- mergeable
- private(list), firstprivate(list)
- shared(list), depend(dependence-type : list)

```
#pragma omp task [clause[[,] clause] ...] new-line  
structured-block
```



Runtime Library Routines

Comprehensive information is available in OpenMP Specification



Examples of Library Routines

- ❑ **omp_set_num_threads(int num_threads)**
 - determines the number of threads
- ❑ **omp_get_num_threads(void)**
 - get the number of threads in the enclosing parallel region
- ❑ **omp_get_max_threads(void)**
 - indicates how many threads could be used
- ❑ **omp_get_thread_num(void)**
 - get the number of the thread (its ID) within the team
- ❑ **omp_get_num_procs(void)**
 - get the number of available processing elements
- ❑ **omp_set_nested(int nested)**
 - enables the nested parallelism



Examples of Lock Routines

- ❑ **OpenMP lock routines may be used for synchronization**
 - affect all tasks that call the routine
 - there are **simple** and **nestable** (may be set multiple times) lock routines
- ❑ **`omp_init_lock(omp_lock_t *lock)`**
 - initializes a simple OpenMP lock in the unlocked state
- ❑ **`omp_set_lock(omp_lock_t *lock)`**
 - suspends the task execution until the lock is available and it is set
- ❑ **`omp_unset_lock(omp_lock_t *lock)`**
 - unlocks the lock
- ❑ **`omp_test_lock(omp_lock_t *lock)`**
 - attempts to set an OpenMP lock without suspending the task execution
- ❑ **`omp_destroy_lock(omp_lock_t *lock)`**
 - deinitializes the lock



Examples of Timing Routines

❑ `omp_get_wtime(void)`

- get the wall clock time in seconds; it is not consistent across all threads

❑ `omp_get_wtick(void)`

- indicates the precision of the timer; number of seconds between clock ticks

```
double begin, end;

begin = omp_get_wtime();
...
end = omp_get_wtime();

printf("Execution time: %f [s]\n", end - begin);
```



Environment Variables

Comprehensive information is available in OpenMP Specification



Examples of Environment Variables

❑ Hello World example

- `export OMP_NUM_THREADS=12`
- determines the number of threads used in parallel region

❑ `OMP_SCHEDULE`

- affects the schedule type and chunk size of loop directives with schedule type runtime
- for instance, `export OMP_SCHEDULE "guided,12"`

❑ Thread affinity

- may improve the locality and memory access
- `OMP_PLACES`: may be threads (hardware threads), cores, sockets (comprises one or more cores)
- `OMP_PROC_BIND`: sets the thread affinity



Example: Matrix Multiplication



Matrix Multiplication

```
// initialization
...
begin = omp_get_wtime();
#pragma omp parallel for shared(a,b,c)
    for (int i = 0; i < dim; ++i) {
        for (int j = 0; j < dim; ++j) {
            for (int k = 0; k < dim; ++k) {
                c[i][j] += a[i][k] * b[k][j];
            }
        }
    }
end = omp_get_wtime();
printf("Execution time: %f [s]\n", end - begin);
...
```



Summary

- ❑ **We have highlighted some features of OpenMP 4.0**
 - more details are provided in the corresponding specification
- ❑ **Various compilers support various versions of OpenMP**
 - OpenMP 4.0 is supported by GCC 4.9 or later version
 - OpenMP 3.0 is supported by GCC 4.4 or later version
 - not everything that is described in specification is supported by a certain compiler
- ❑ **Currently in our lab is installed**
 - gcc version 5.4.0

